

Chapter 8
Passing Parameters to Subroutines
via Registers

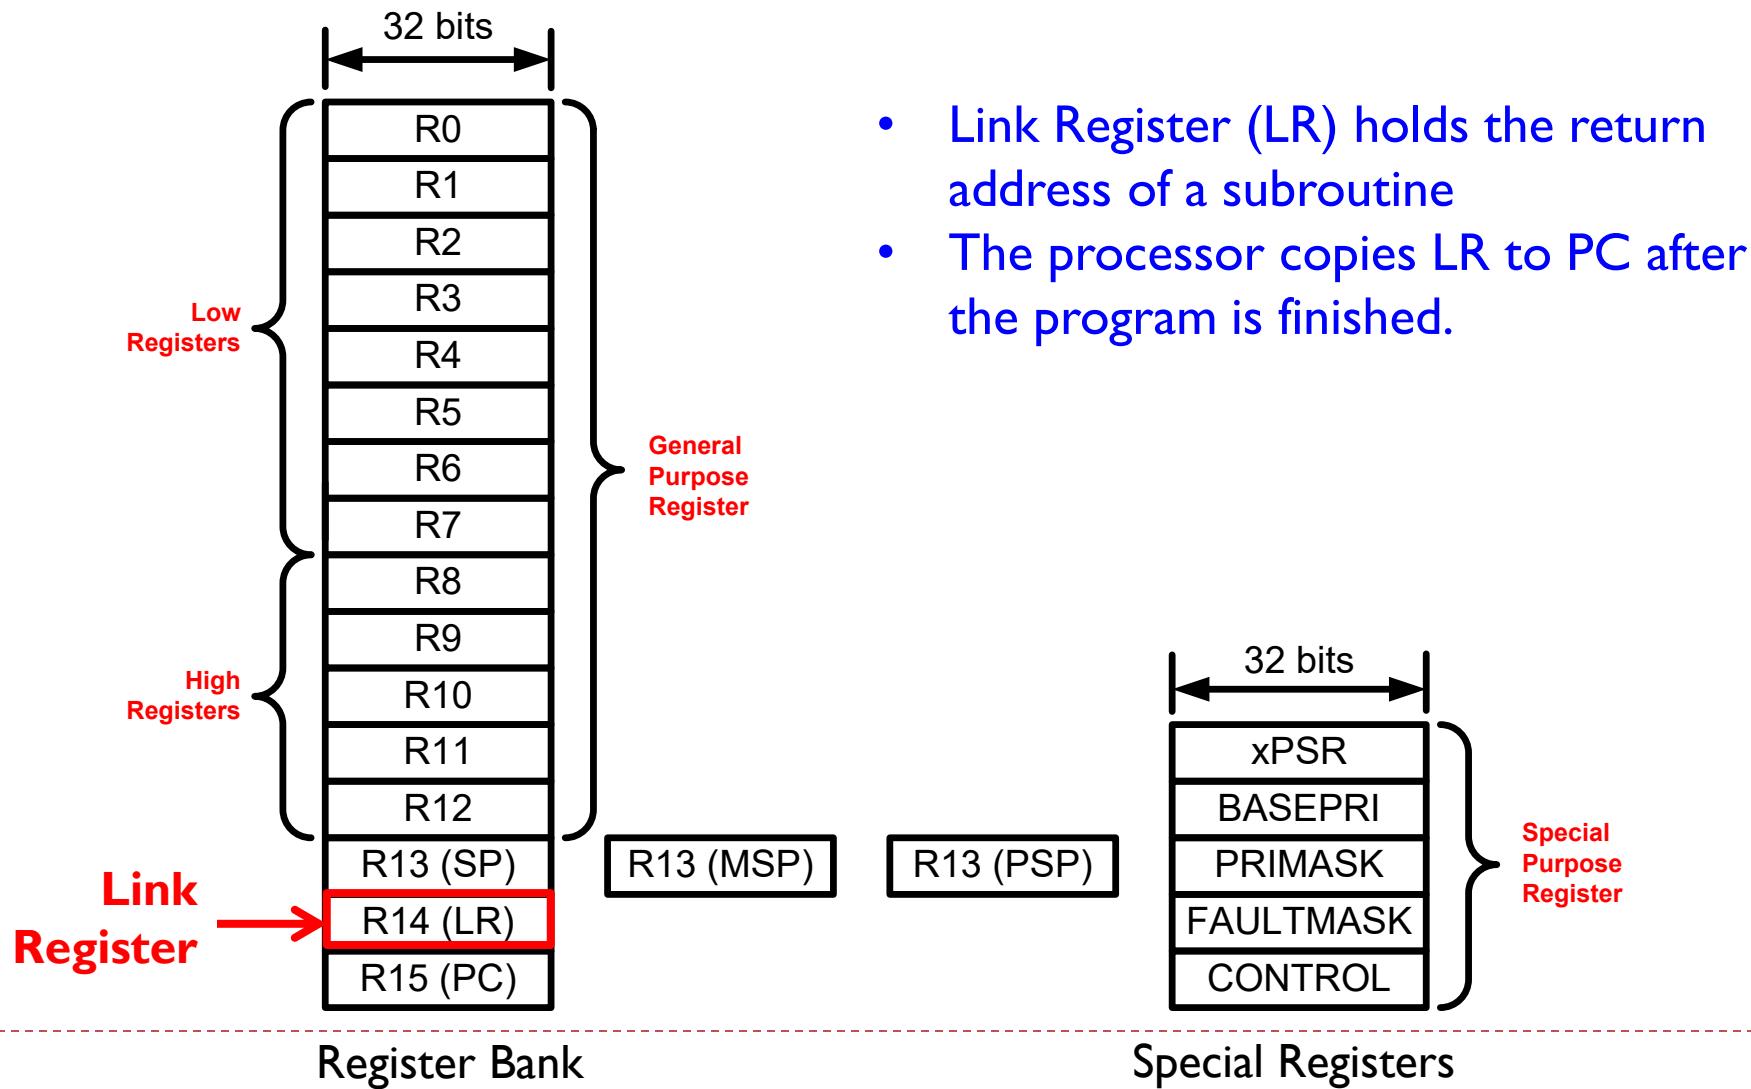
Z. Gu

Fall 2025

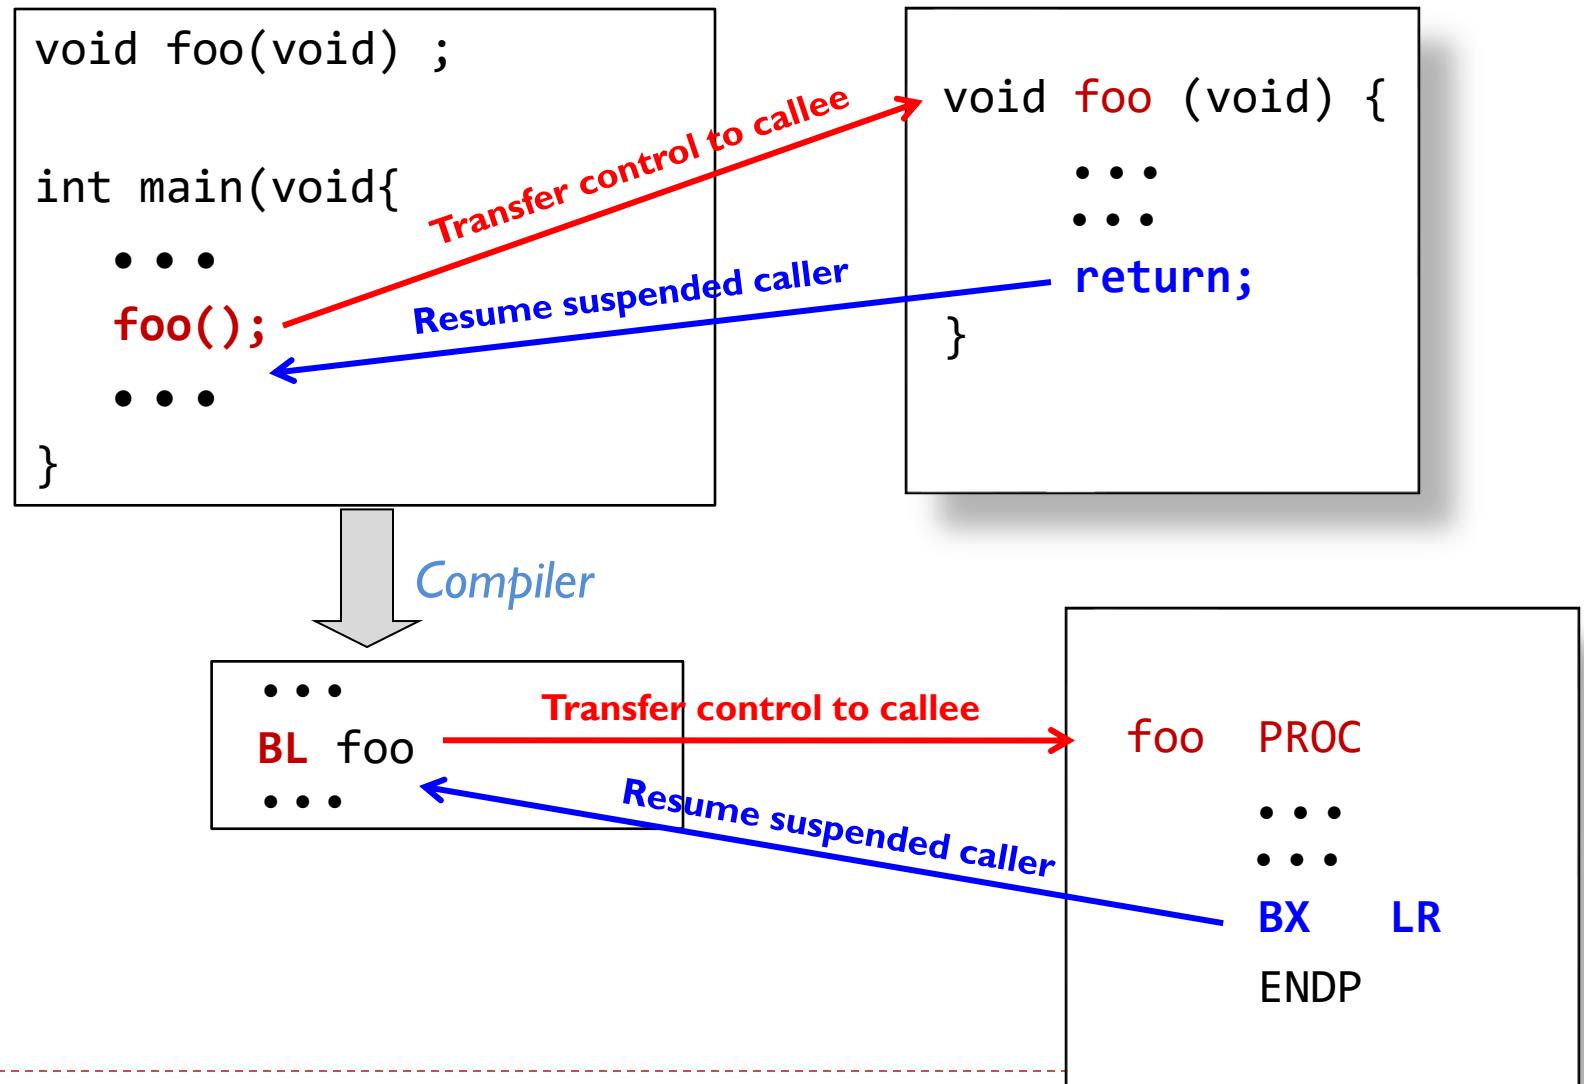
Overview

- ▶ How to call a subroutine?
- ▶ How to return the control back to the caller?
- ▶ How to pass arguments into a subroutine?
- ▶ How to return a value in a subroutine?
- ▶ How to preserve the running environment for the caller?

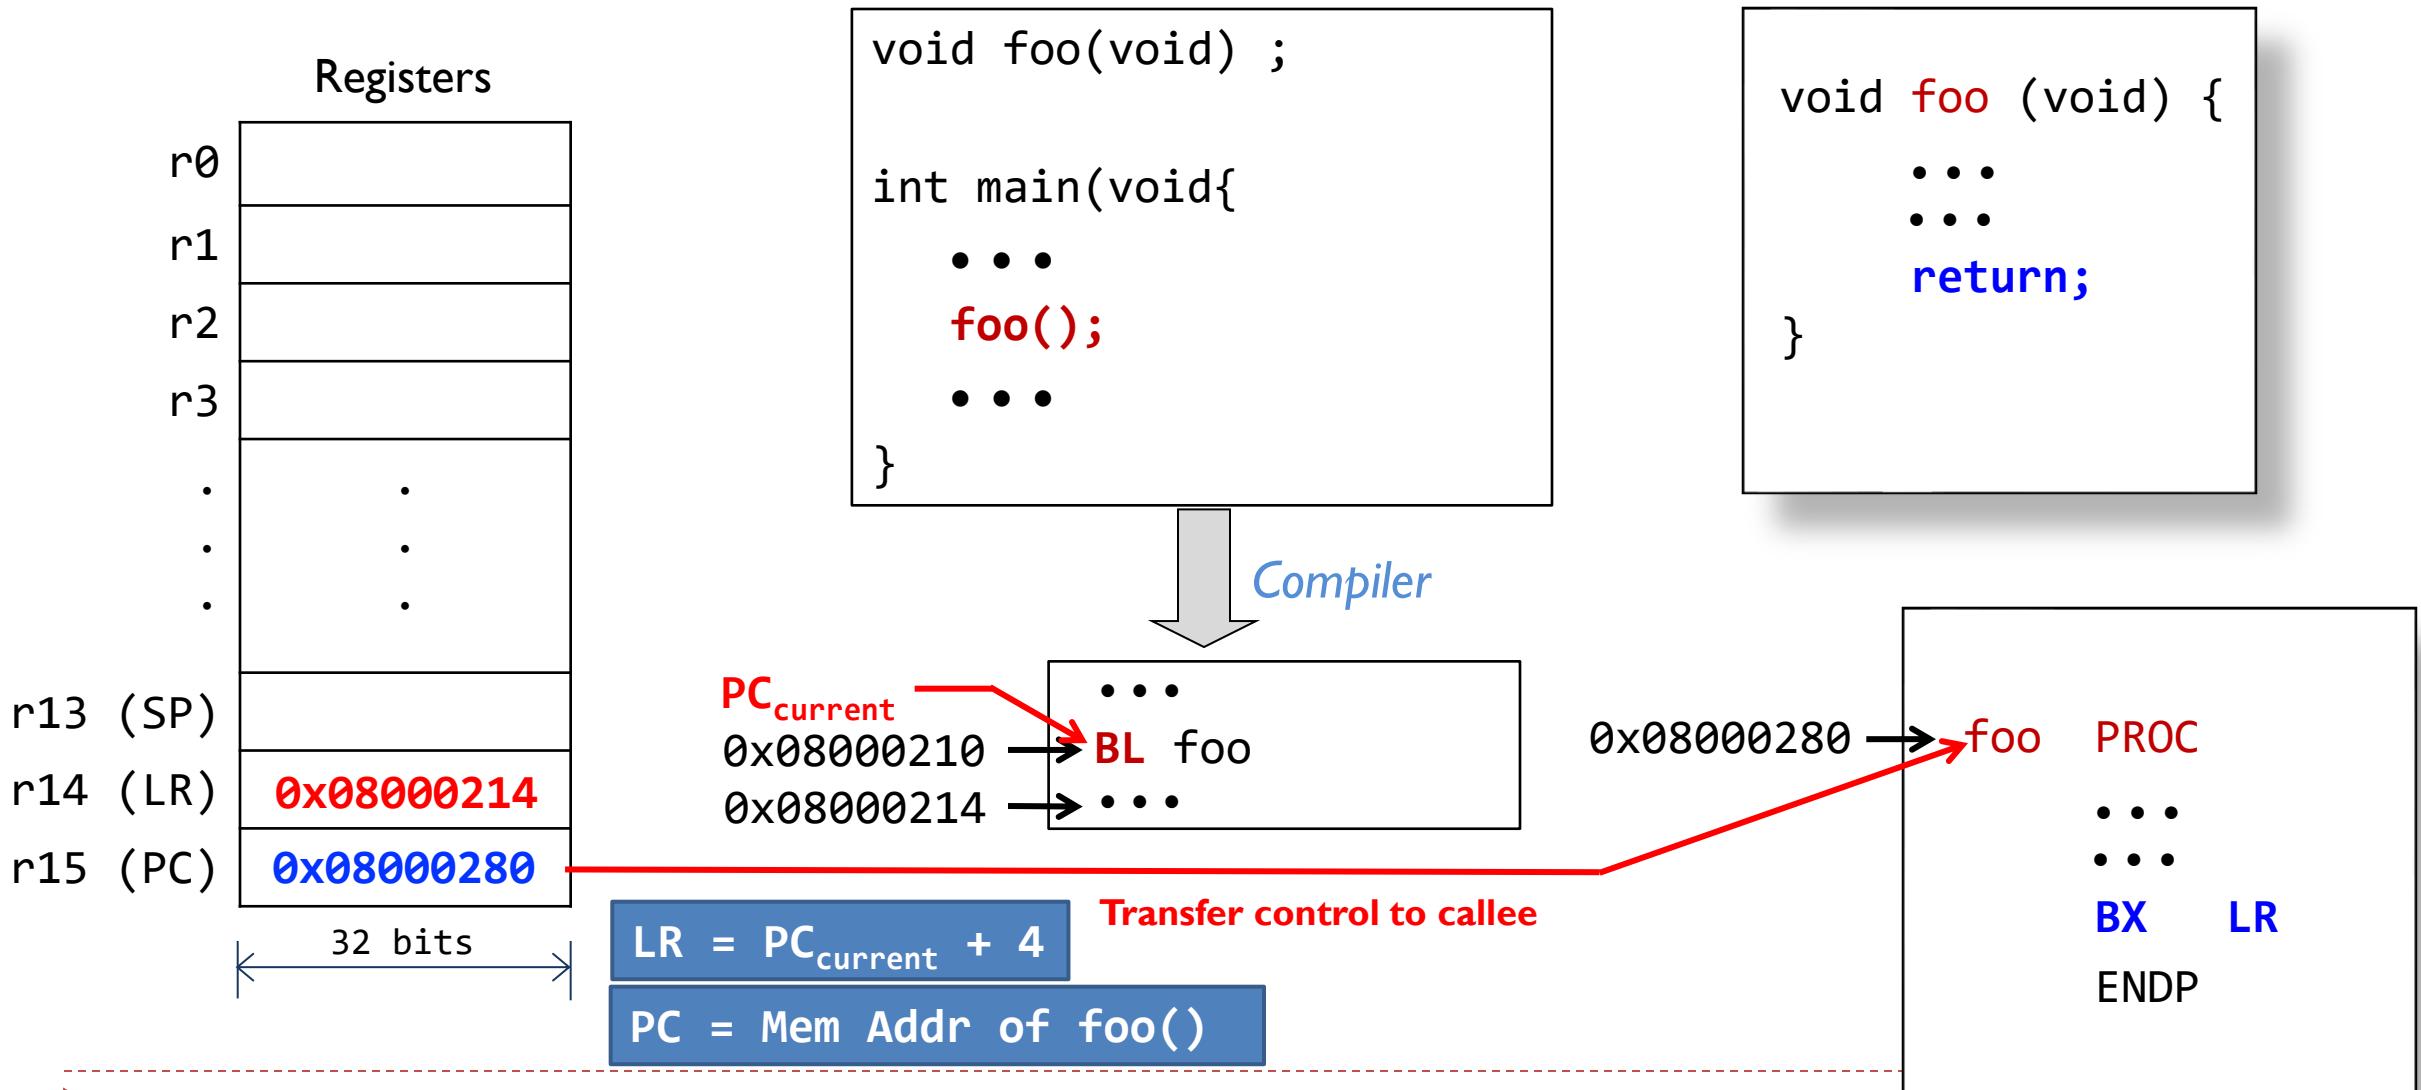
Link Register (LR)



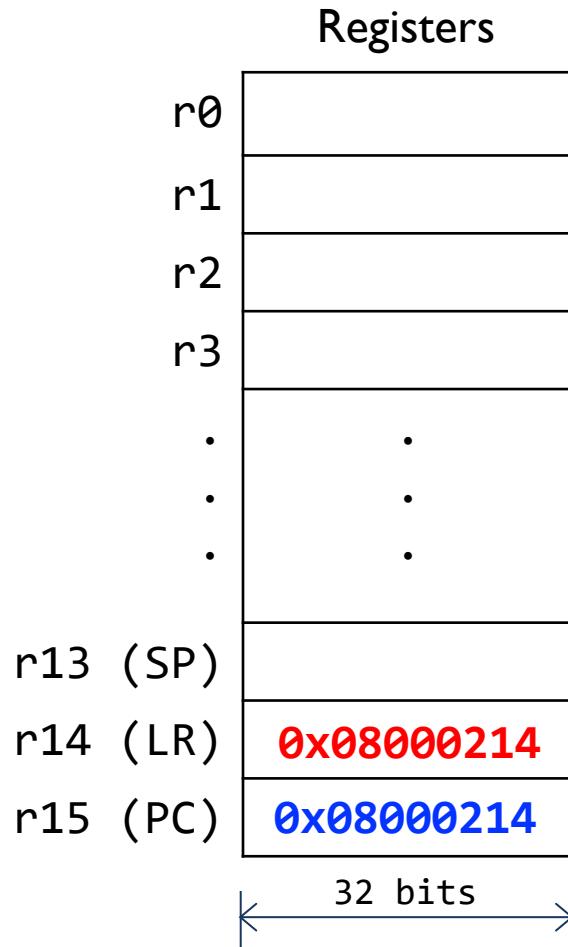
Link Register (LR)



Link Register (LR)



Link Register (LR)



```
void foo(void);  
  
int main(void{  
    ...  
    foo();  
    ...  
})
```

Compiler
↓

```
...  
0x08000210 → BL foo  
0x08000214 → ...
```

Resume suspended caller

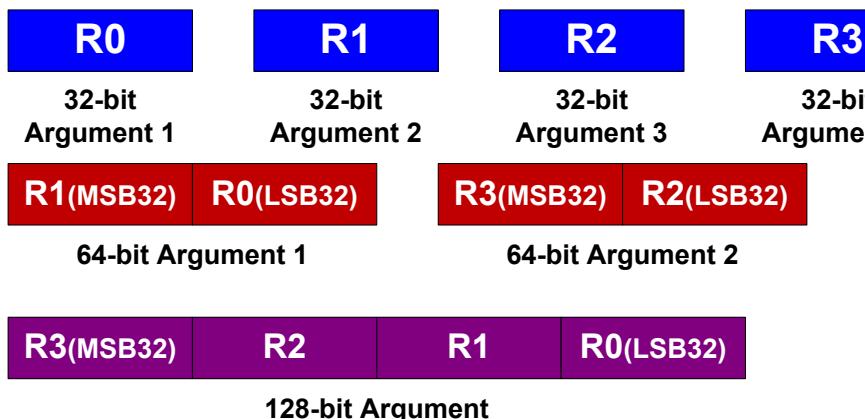
```
void foo (void) {  
    ...  
    return;  
}  
  
0x08000280 → foo PROC  
...  
...  
BX LR  
ENDP
```

PC_{current}

Procedure Call Standard

- ▶ What is it?
 - ▶ Contract between a calling subroutine (caller) and a called subroutine (callee)
- ▶ Why need it?
 - ▶ Allows subroutines to be separately written, compiled, assembled but work together
 - ▶ Allows C program calls an assembly function, or vice versa
- ▶ This talk focuses on
 - ▶ How to pass arguments to a subroutine?
 - ▶ How to return a result from a subroutine?

Passing Arguments and Returning Value



Extra arguments are pushed to the stack by the caller. The caller is responsible to pop them out of the stack after the subroutine returns.

- ▶ Each argument with size \leq 32 bits, e.g., 8-bit char, or 16-bit short, or 32-bit int, is passed in a 32-bit register.
 - ▶ Cannot pack multiple arguments into one register.
- ▶ The subroutine can take arguments larger than 32 bits. For example, a double-word variable, such as 64-bit long, is passed in two consecutive registers (e.g. R0 and R1, or R2 and R3). A 128-bit variable is passed in four consecutive registers.
 - ▶ `int64_t add_64(int64_t a, int64_t b)`
 - ▶ R0 and R1 are used to store the variable `a`
- ▶ The return result is stored in registers (R0-R3), depending on the size of the return variable. If it is less than 32 bits, it is stored in R0. If it is a double-word sized variable, such as `long long` or `double` variables in C, it is stored in R0 and R1.
 - ▶ `int128_t multiply_64(int64_t a, int64_t b)`
 - ▶ R0, R1, R2, and R3 are used to store the result

The diagram shows the return value sizes corresponding to the argument sizes. A brown box labeled "32-bit Return Value" is aligned with the R0 register. A blue box labeled "64-bit Return Value" is aligned with the R1(MSB32) and R0(LSB32) pair. A purple box labeled "128-bit Return Value" is aligned with the R3(MSB32), R2, R1, and R0(LSB32) sequence.

R0	R1(MSB32)	R0(LSB32)	R3(MSB32)	R2	R1	R0(LSB32)
32-bit Return Value	64-bit Return Value		128-bit Return Value			

Passing Arguments Examples

Registers

R0	R1	R2	R3
----	----	----	----

Stack in Memory

--	--	--	--	--	--	--

foo (int i0, int i1, int i2, int i3)

Registers

i0	i1	i2	i3
----	----	----	----

Stack in Memory

--	--	--	--	--	--	--

foo (int i0, char a1, double D)

Registers

i0	a1	D
----	----	---

Stack in Memory

--	--	--	--	--	--	--

foo (int i0, int i1, double D, int i2, int i3)

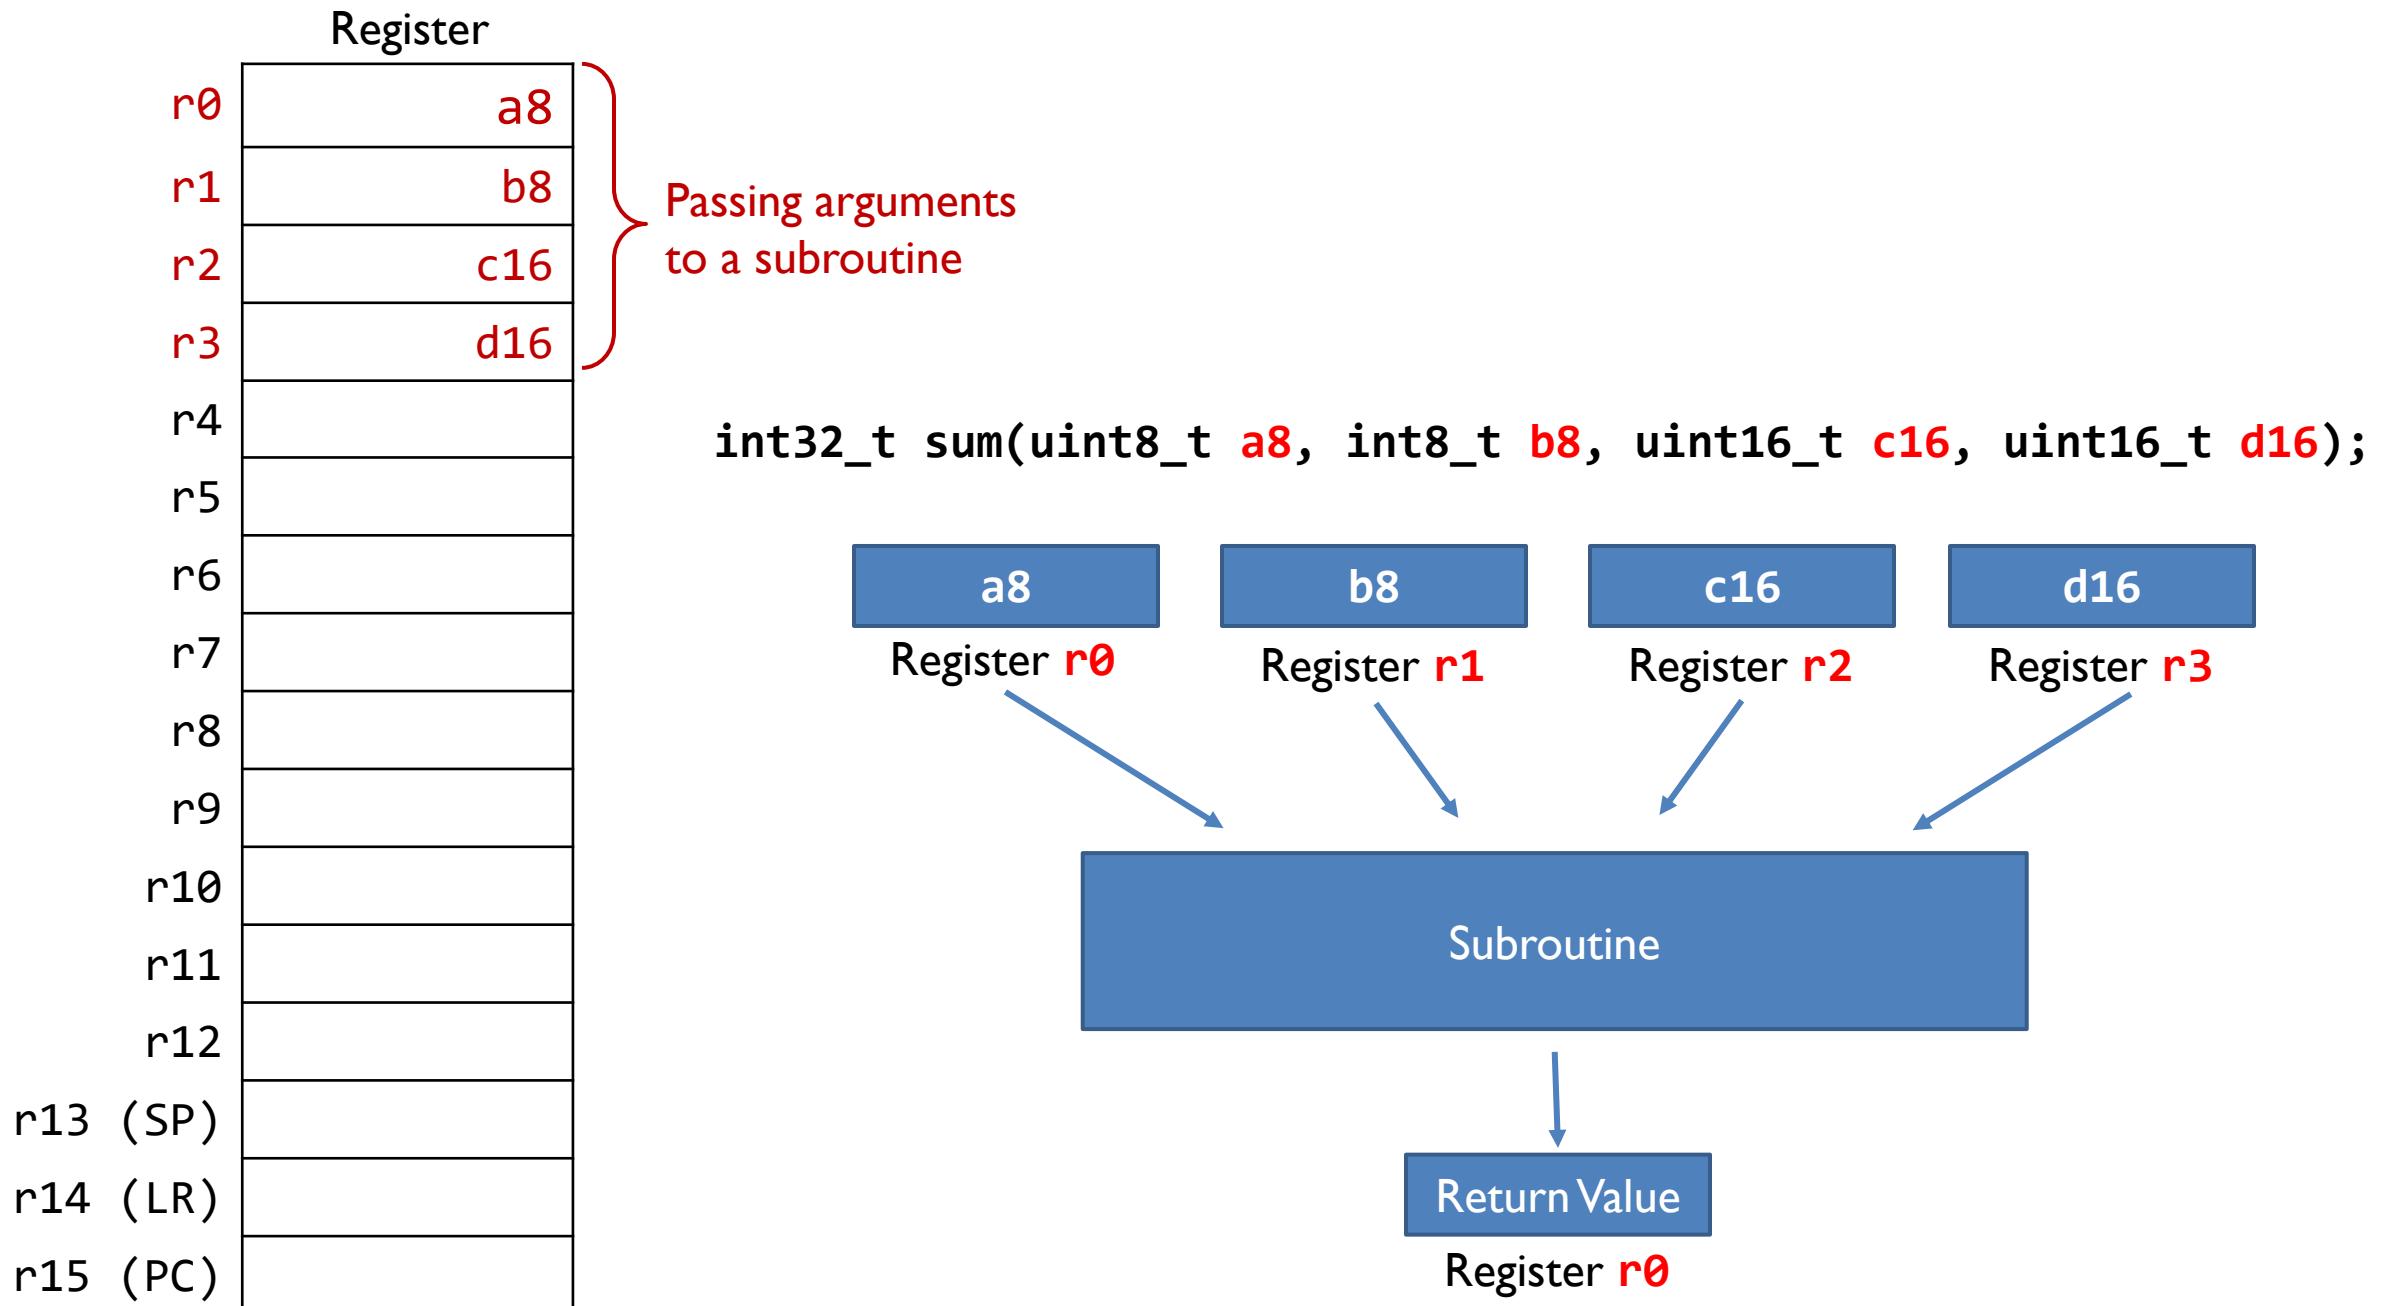
Registers

i0	i1	D
----	----	---

Stack in Memory

i2	i3					
----	----	--	--	--	--	--

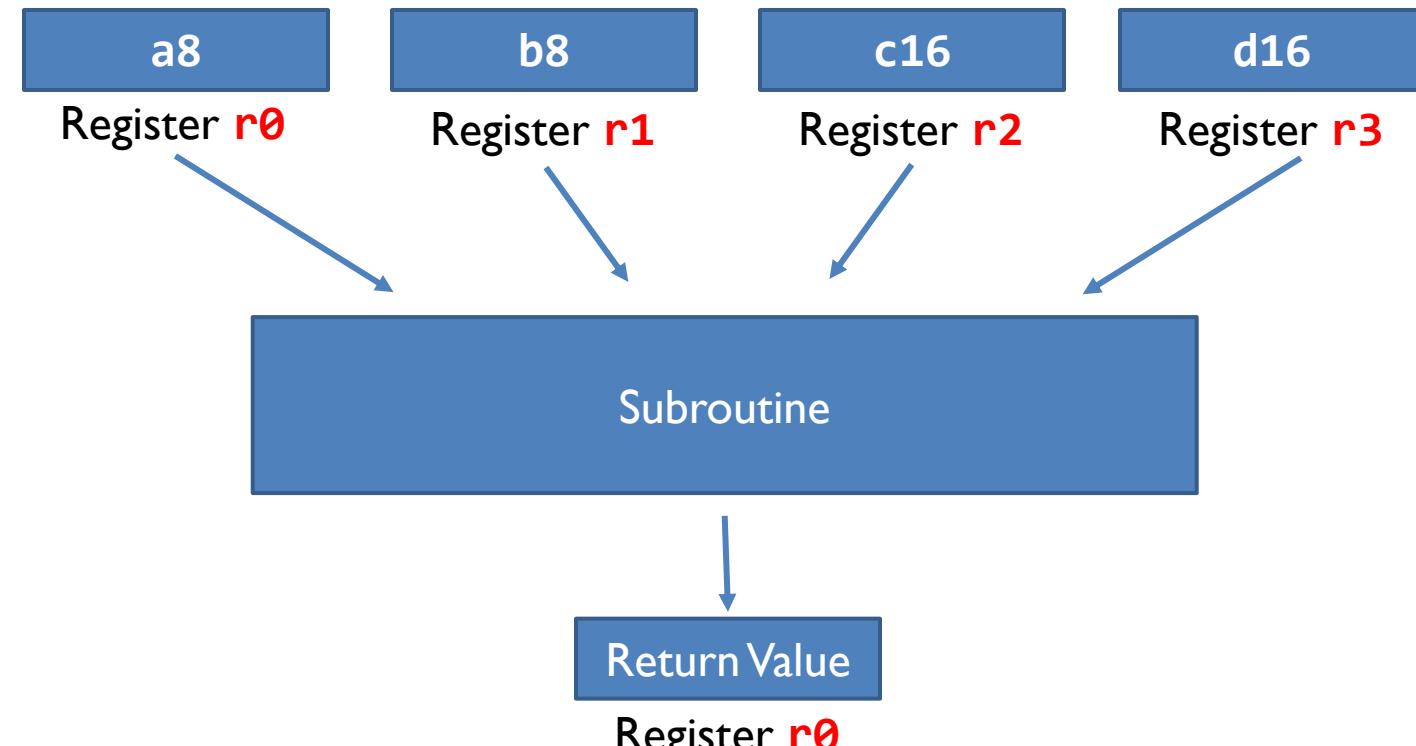
Caller passes arguments i0, i1, D in registers R0-R3 directly; pushes additional arguments i2 and i3 on the stack before subroutine call



Register	
r0	sum
r1	
r2	
r3	
r4	
r5	
r6	
r7	
r8	
r9	
r10	
r11	
r12	
r13 (SP)	
r14 (LR)	
r15 (PC)	

Return the sum in register r0

```
int32_t sum(uint8_t a8, int8_t b8, uint16_t c16, uint16_t d16);
```



Passing Arguments and Returning Value

```
int32_t sum(uint8_t a8, int8_t b8, uint16_t c16, uint16_t d16);  
  
s = sum(1, 2, 3, 4);
```

Caller

```
MOVS r0, #1 ; a8  
MOVS r1, #2 ; b8  
MOVS r2, #3 ; c16  
MOVS r3, #4 ; d16  
BL sum
```

Callee

```
sum PROC  
    ADD r0, r0, r1 ; a8 + b8  
    ADD r0, r0, r2 ; add c16  
    ADD r0, r0, r3 ; add d16  
    BX LR  
ENDP
```

Returning Value

```
uint32_t s32;  
  
uint32_t sum(uint8_t a8, uint8_t b8, uint16_t c16, uint16_t d16);  
  
s32 = sum(1, 2, 3, 4) + 100;
```

```
MOVS r0, #1 ; 1st argument a8  
MOVS r1, #2 ; 2nd argument b8  
MOVS r2, #3 ; 3rd argument c16  
MOVS r3, #4 ; 4th argument d16  
BL sum ; result is returned in r0  
ADD r0, r0, #100  
LDR r4, =s32 ; Get memory address of s32  
STR r0, [r4] ; Save returned result to s32
```

uint32_t s32 is declared as a C global variable, so the compiler/linker creates storage for it in .data or .bss depending on initialization. You can simply reference it by name in assembly as a label.

Register	
r0	a[31:0]
r1	a[63:32]
r2	b[31:0]
r3	b[63:32]
r4	
r5	
r6	
r7	
r8	
r9	
r10	
r11	
r12	
r13 (SP)	
r14 (LR)	
r15 (PC)	

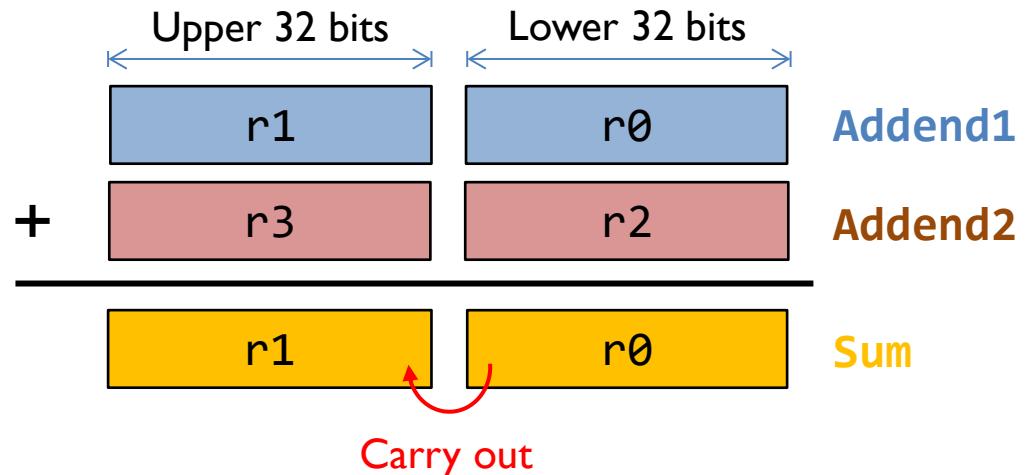
Passing arguments
to a subroutine

`int64_t sum(int64_t a, int64_t b);`

Callee

sum PROC

```
ADDS r0, r2, r0 ; Adding lower 32 bits
ADC r1, r3, r1 ; Adding upper 32 bits
BX LR           ; Return in r1:r0
ENDP
```



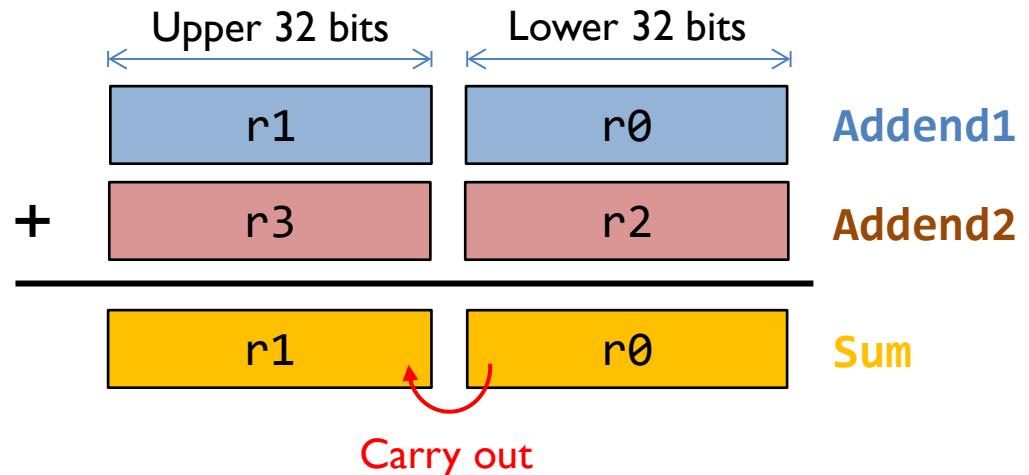
Register	
r0	sum[31:0]
r1	sum[63:32]
r2	
r3	
r4	
r5	
r6	
r7	
r8	
r9	
r10	
r11	
r12	
r13 (SP)	
r14 (LR)	
r15 (PC)	

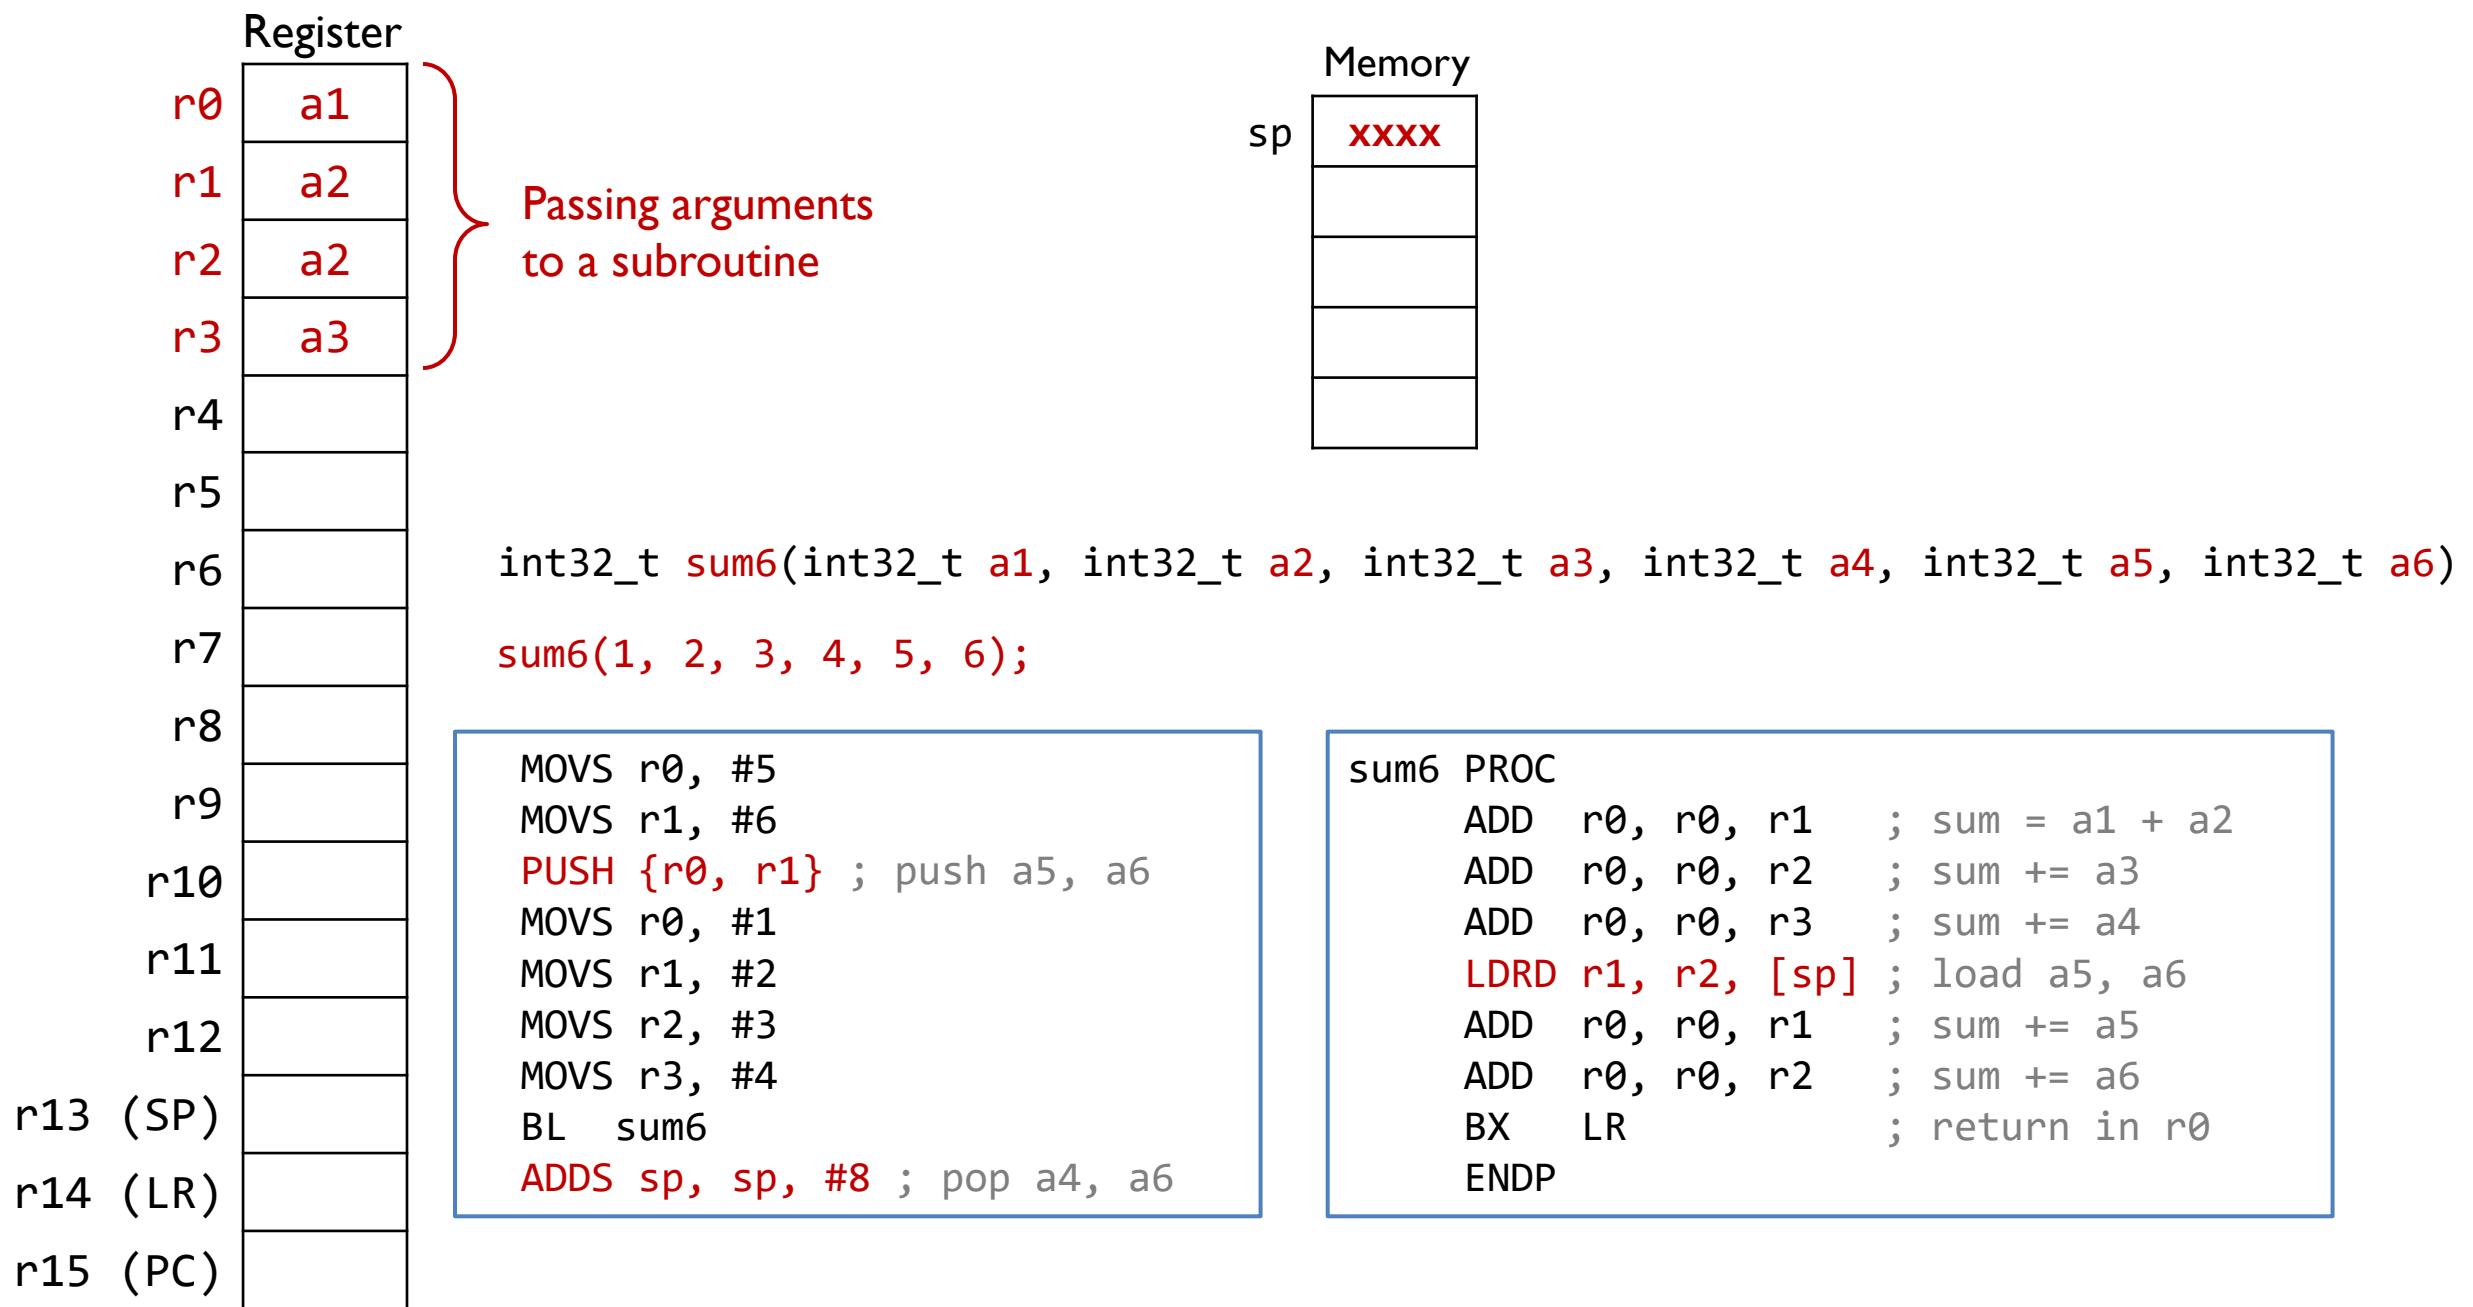
Returning sum
from a subroutine

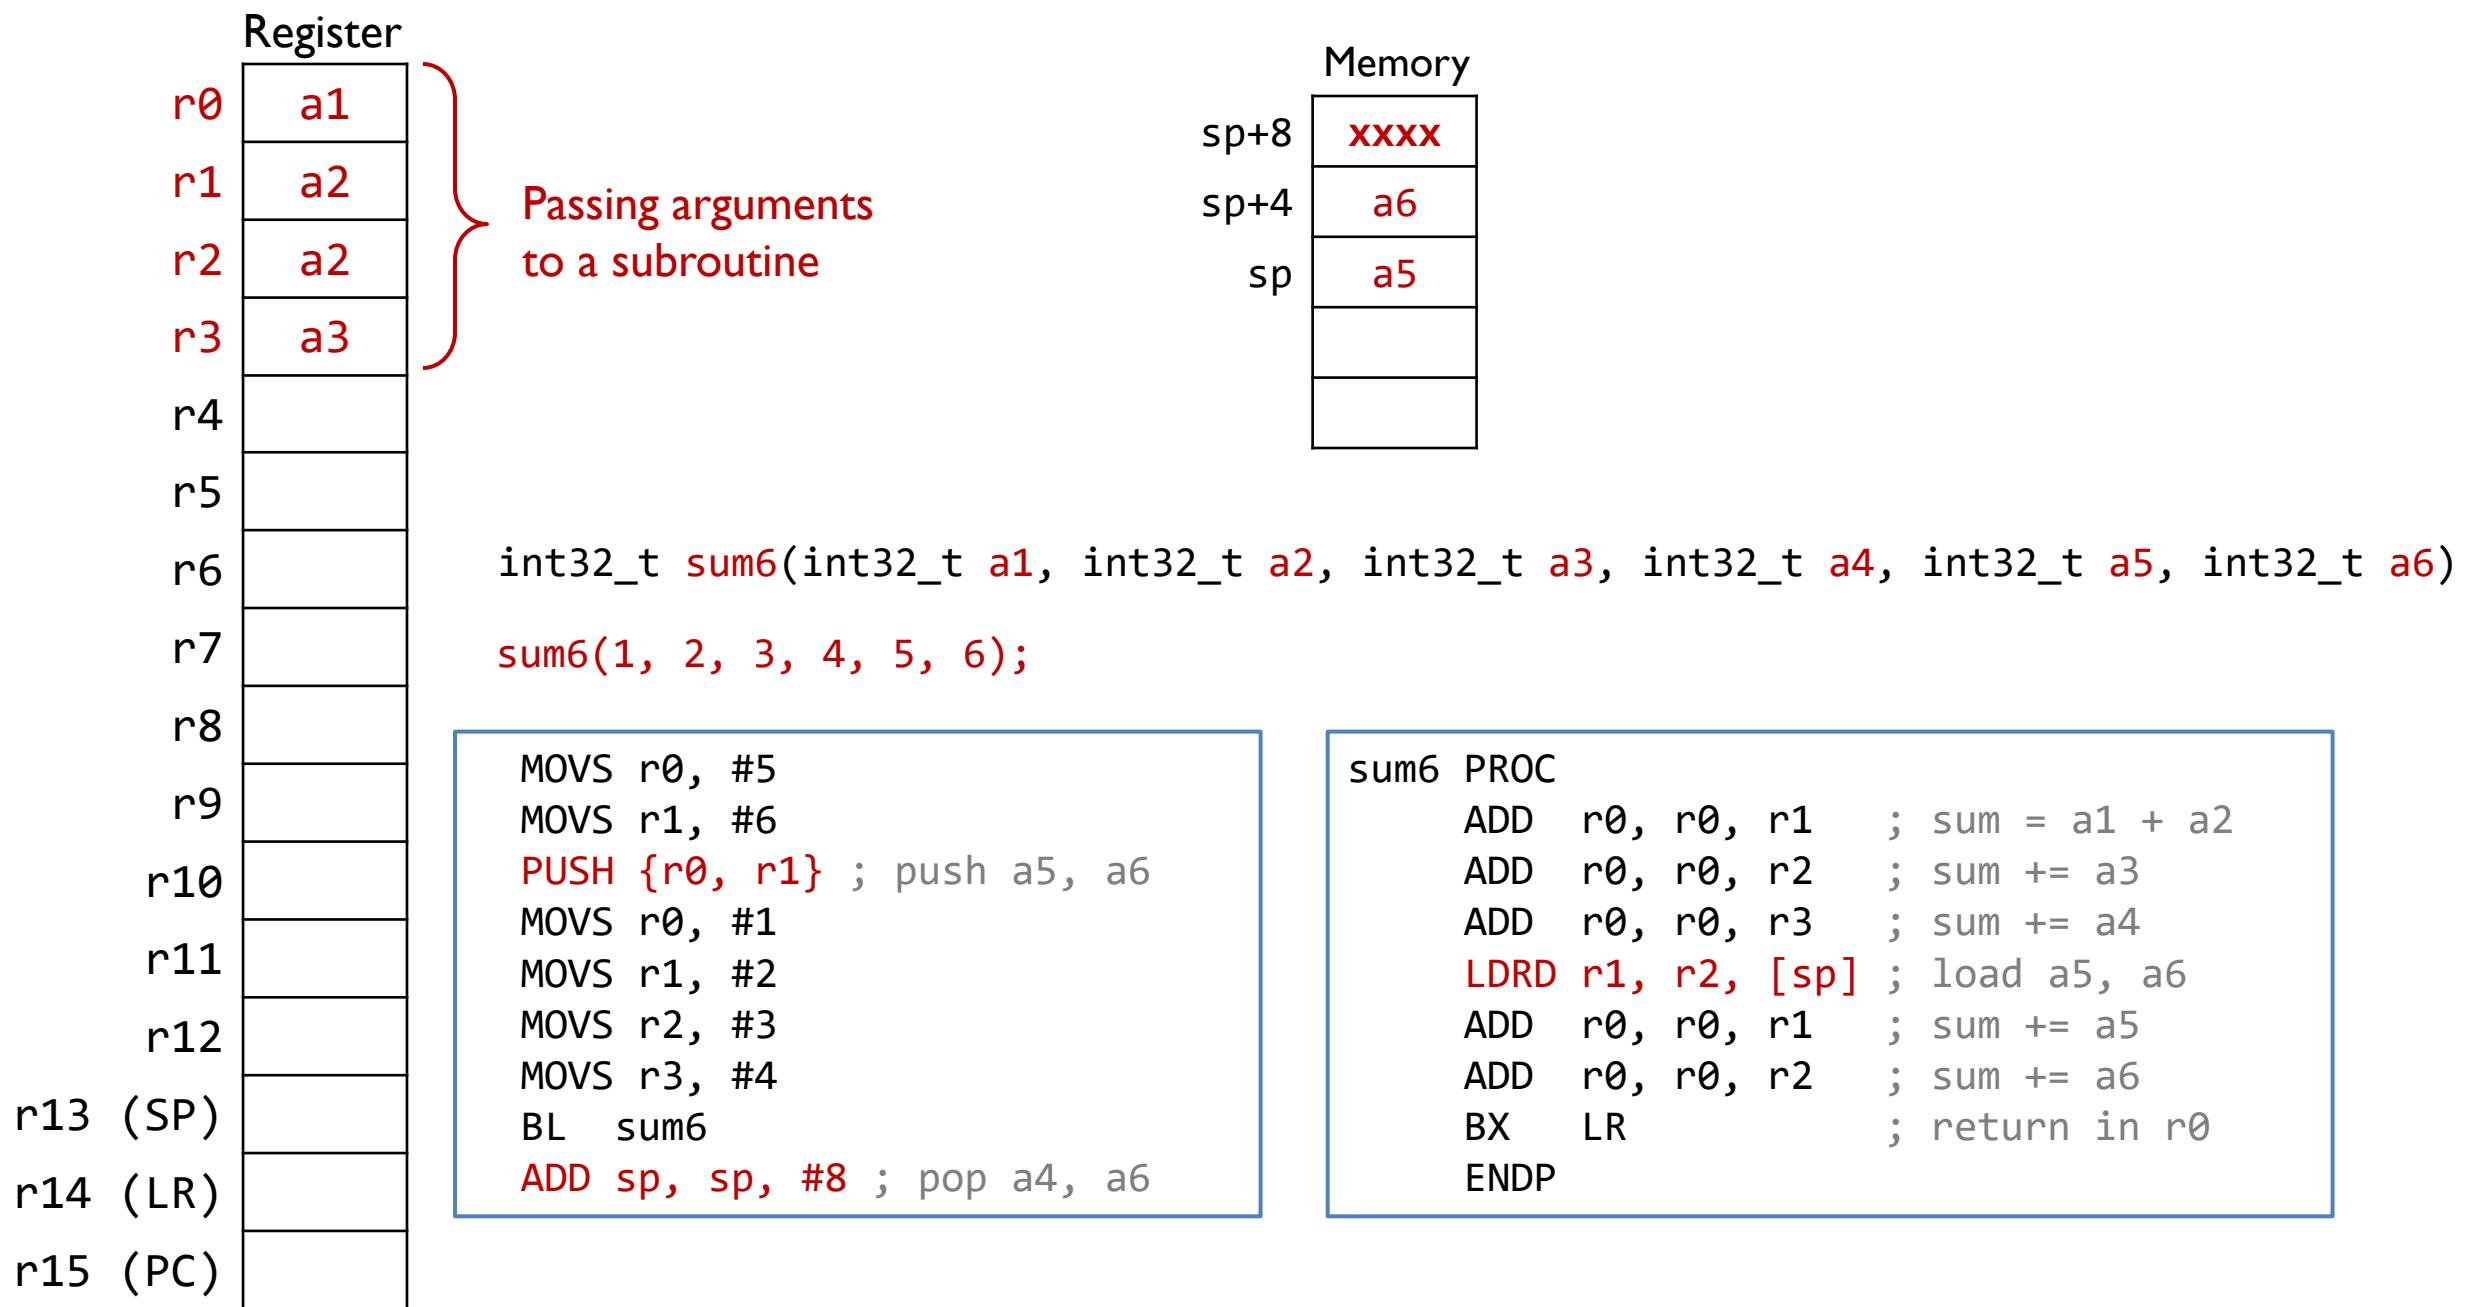
`int64_t sum(int64_t a, int64_t b);`

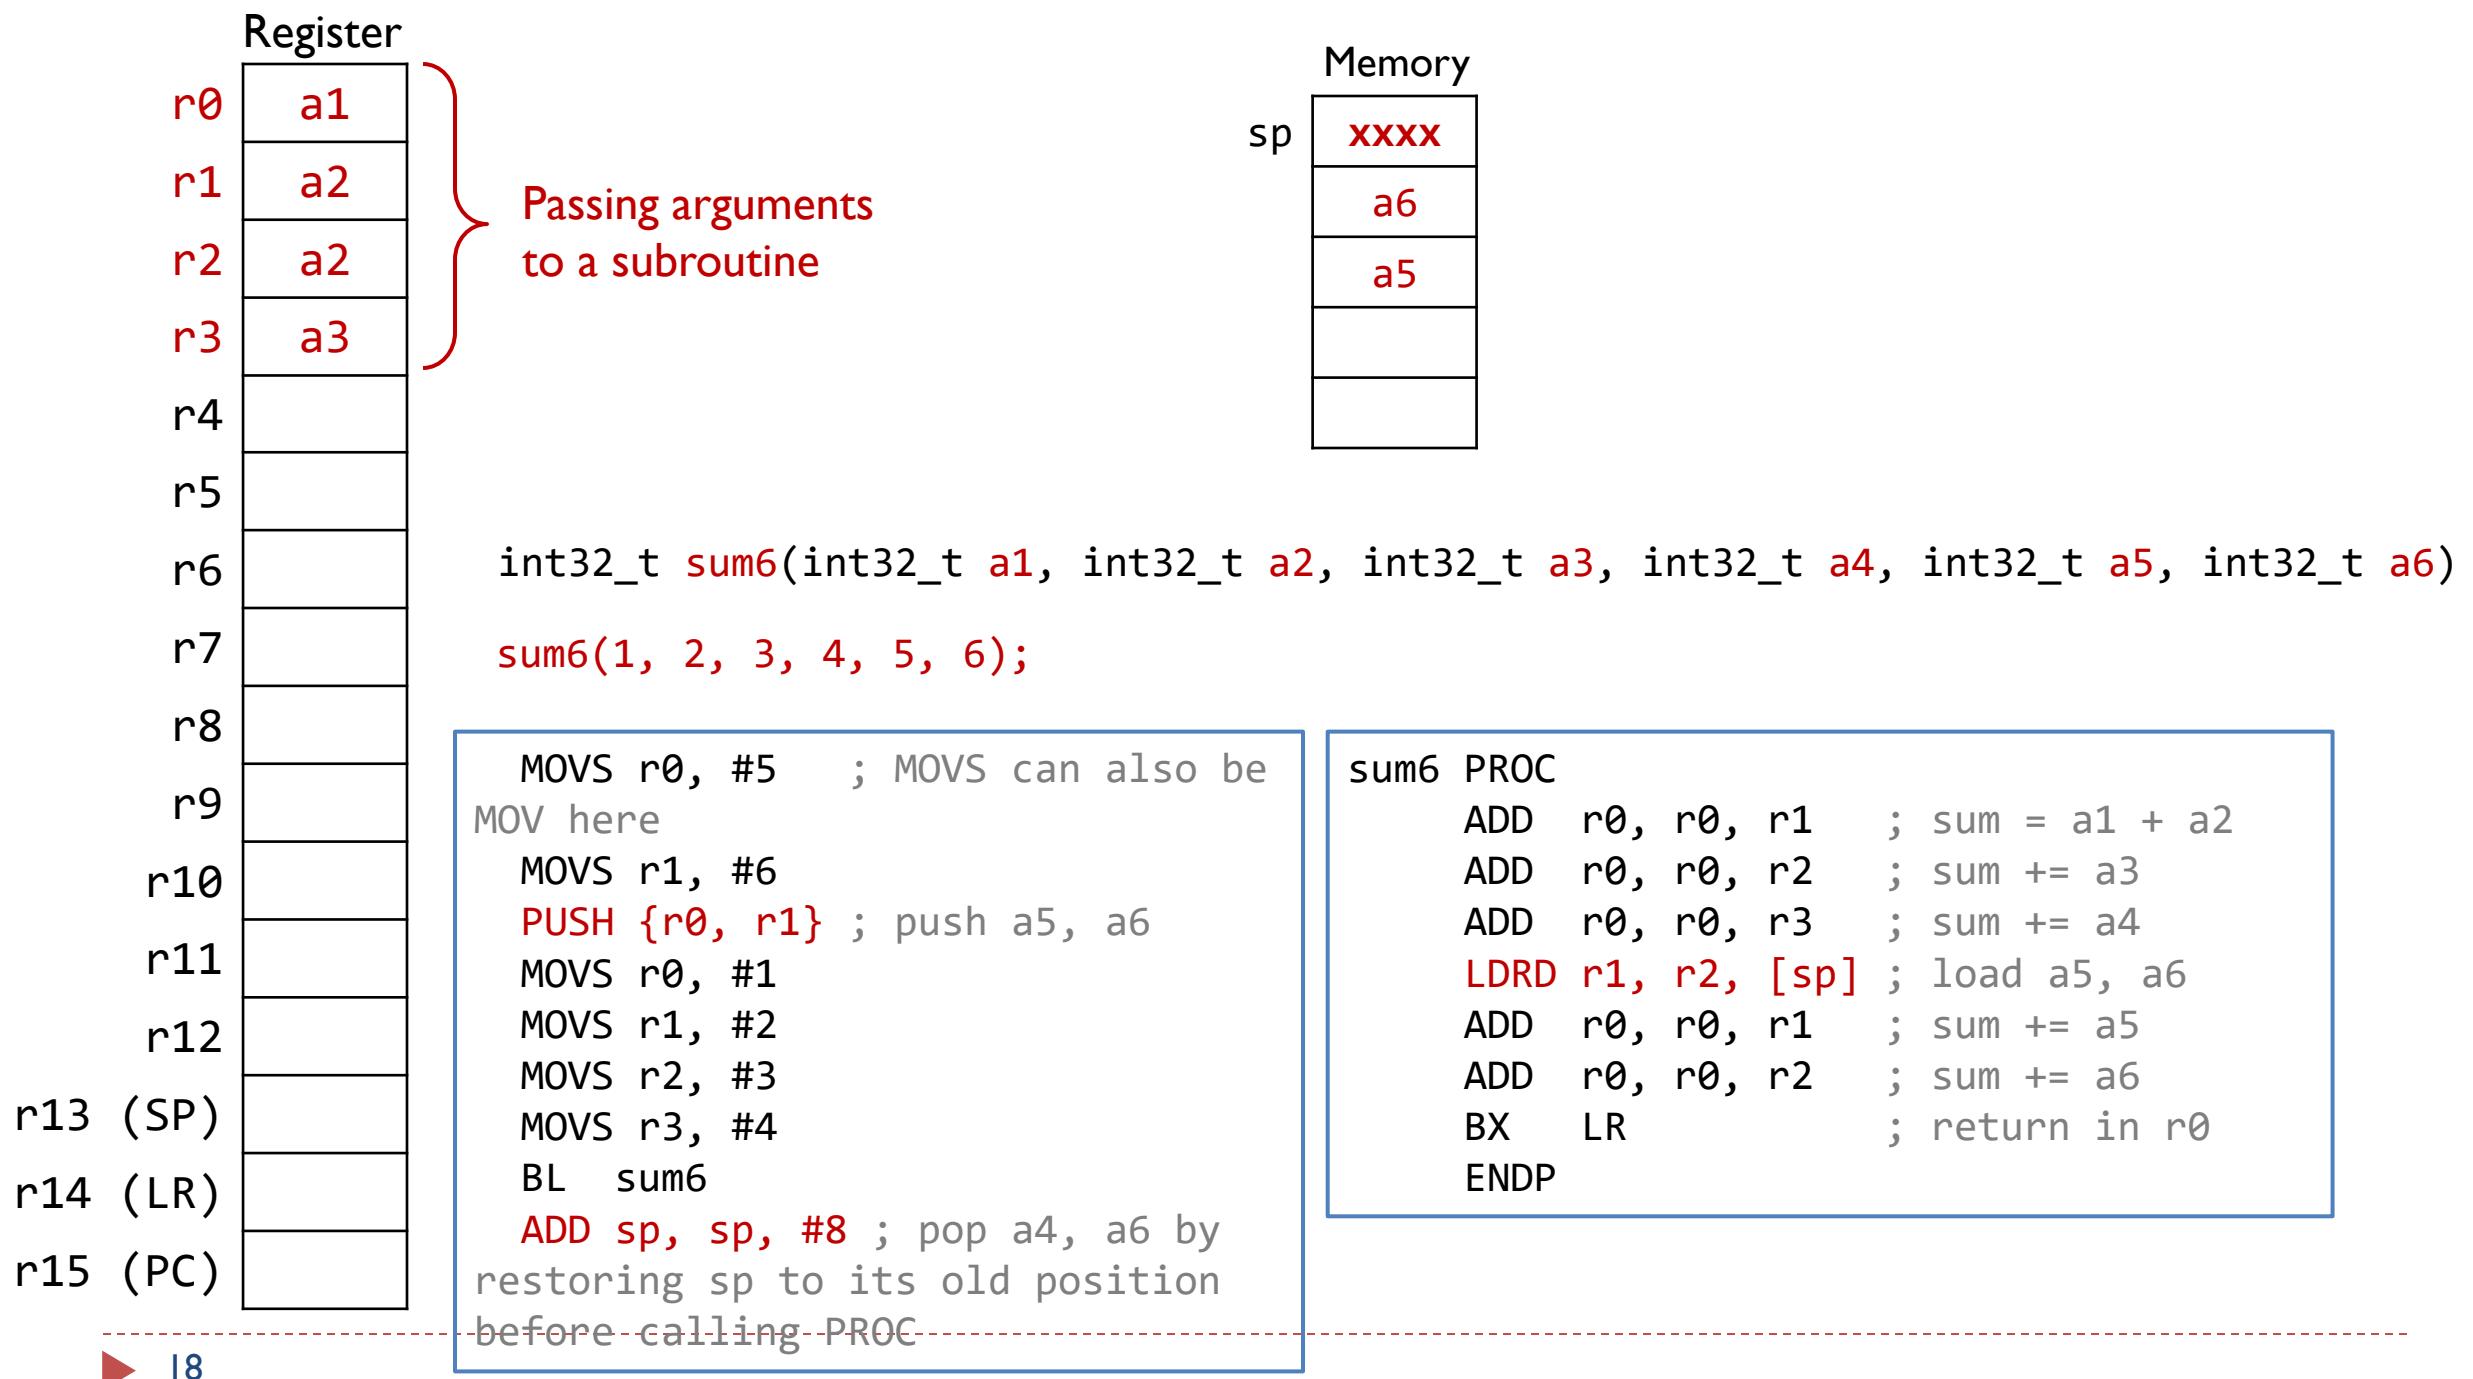
Callee

```
sum PROC
    ADDS r0, r2, r0 ; Adding lower 32 bits
    ADC r1, r3, r1 ; Adding upper 32 bits
    BX LR           ; Return in r1:r0
ENDP
```









Explanations

Step	r0	r1	r2	r3	Stack (top first)	Description
Before call	5	6	—	—	—	a5, a6 prepared
After PUSH	—	—	—	—	a5=5, a6=6	pushed to stack
Set r0–r3	1	2	3	4	a5=5, a6=6	a1–a4 in regs
Inside sum6	1	2	3	4	a5=5, a6=6	entry point
LDRD loads	—	5	6	—	a5=5, a6=6	from stack
Return	21	—	—	—	a5=5, a6=6	1+2+3+4+5+6 = 21

- ▶ First 4 parameters → registers r0–r3. Extra parameters → pushed onto the stack.
- ▶ LDRD (Load Register Double) fetches two words (a5, a6) efficiently.
- ▶ r0 always holds the return value.
- ▶ Caller cleans up the stack after the function (ADDS sp, sp, #8).

Quiz

- ▶ Can I use POP to replace ADD?

```
MOVS r0, #5
MOVS r1, #6
PUSH {r0, r1} ; push a5, a6
MOVS r0, #1
MOVS r1, #2
MOVS r2, #3
MOVS r3, #4
BL sum6
POP {r0, r1}
; ADD sp, sp, #8
```

Quiz ANS

- ▶ POP {r0, r1} is equivalent to LDMIA sp!, {r0, r1}
 - ▶ Load [sp] into r0
 - ▶ Load [sp + 4] into r1
 - ▶ Then increment sp by 8
- ▶ This is wrong, because it would overwrite registers r0 and r1 with garbage (the old arguments you pushed), whereas r0 should contain the return value
 - ▶ After returning from sum6, you do not need those values (a5, a6) anymore. You just want to discard them, so ADD moves stack pointer up by 8 bytes and discards data.

```
MOVS r0, #5
MOVS r1, #6
PUSH {r0, r1} ; push a5, a6
MOVS r0, #1
MOVS r1, #2
MOVS r2, #3
MOVS r3, #4
BL sum6
POP {r0, r1}
; ADD sp, sp, #8
```

Calling Assembly Subroutine in C

- ▶ If your assembly code follows the procedure call standard, a C code can call an assembly subroutine, and vice versa.

```
extern int32_t sum3(int32_t a1, int32_t a2, int32_t a3);

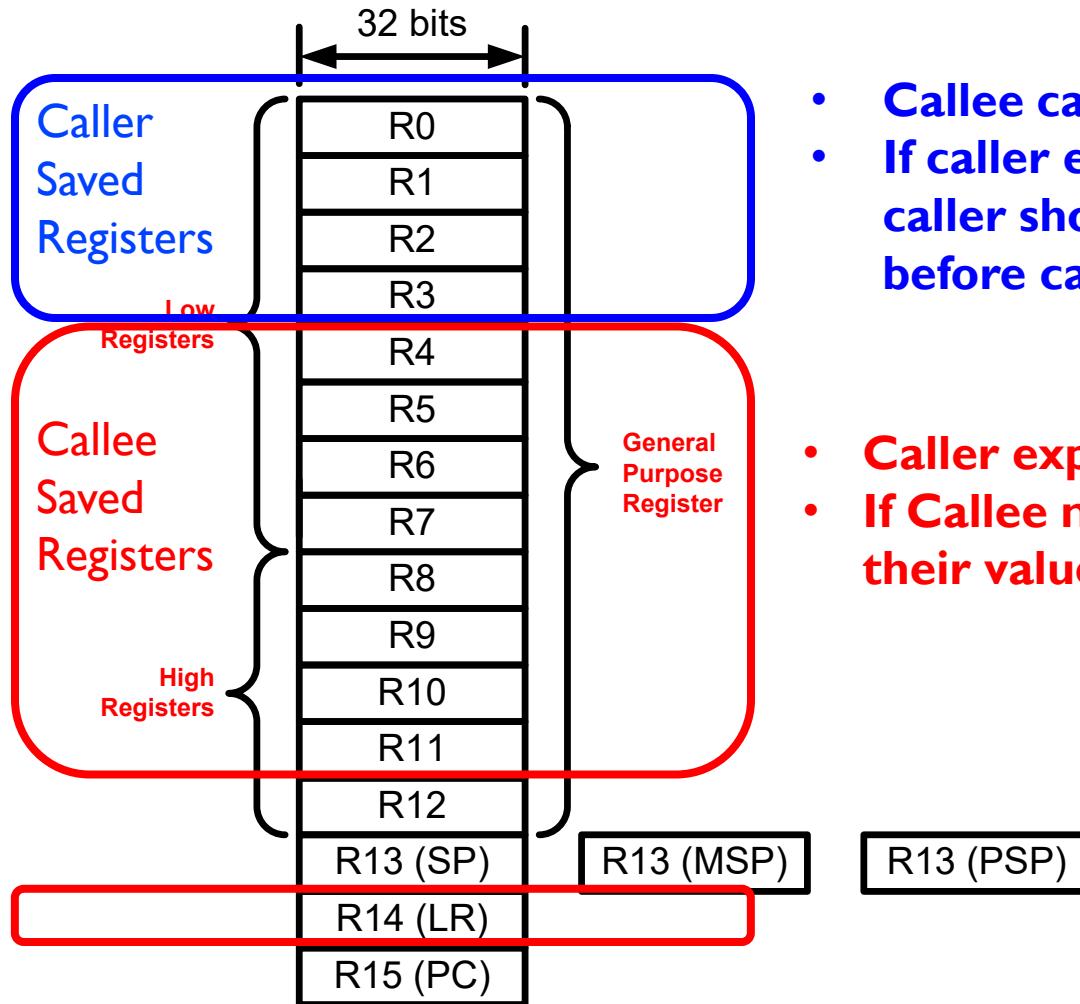
int main(void){
    int32_t s
    ...
    s = sum3(-1, -2, -3) + sum3(4, 5, 6);
    ...
}
```

```
sum3 PROC
    EXPORT sum3
    ADD r0, r0, r1    ; sum = a1 + a2
    ADD r0, r0, r2    ; sum += a3
    BX LR             ; return in r0
ENDP
```

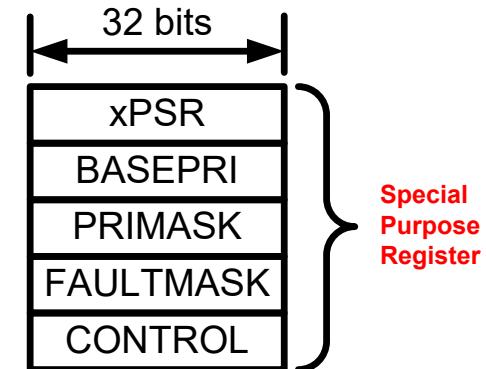
ARM Procedure Call Standard

Register	Usage	Subroutine Preserved	Notes
r0	Argument 1 and return value	No	If return has 64 bits, then r0:r1 hold it. If argument 1 has 64 bits, r0:r1 hold it.
r1	Argument 2	No	
r2	Argument 3	No	If the return has 128 bits, r0-r3 hold it.
r3	Argument 4	No	If more than 4 arguments, use the stack
r4	General-purpose V1	Yes	Variable register 1 holds a local variable.
r5	General-purpose V2	Yes	Variable register 2 holds a local variable.
r6	General-purpose V3	Yes	Variable register 3 holds a local variable.
r7	General-purpose V4	Yes	Variable register 4 holds a local variable.
r8	General-purpose V5	Yes	Variable register 5 holds a local variable.
r9	Platform specific/V6	Yes/No	Usage is platform-dependent.
r10	General-purpose V7	Yes	Variable register 7 holds a local variable.
r11	General-purpose V8	Yes	Variable register 8 holds a local variable.
r12 (IP)	Intra-procedure-call register	No	It holds intermediate values between a procedure and the sub-procedure it calls.
r13 (SP)	Stack pointer	Yes	SP has to be the same after a subroutine has completed.
r14 (LR)	Link register	No	LR does not have to contain the same value after a subroutine has completed.
r15 (PC)	Program counter	N/A	Do not directly change PC

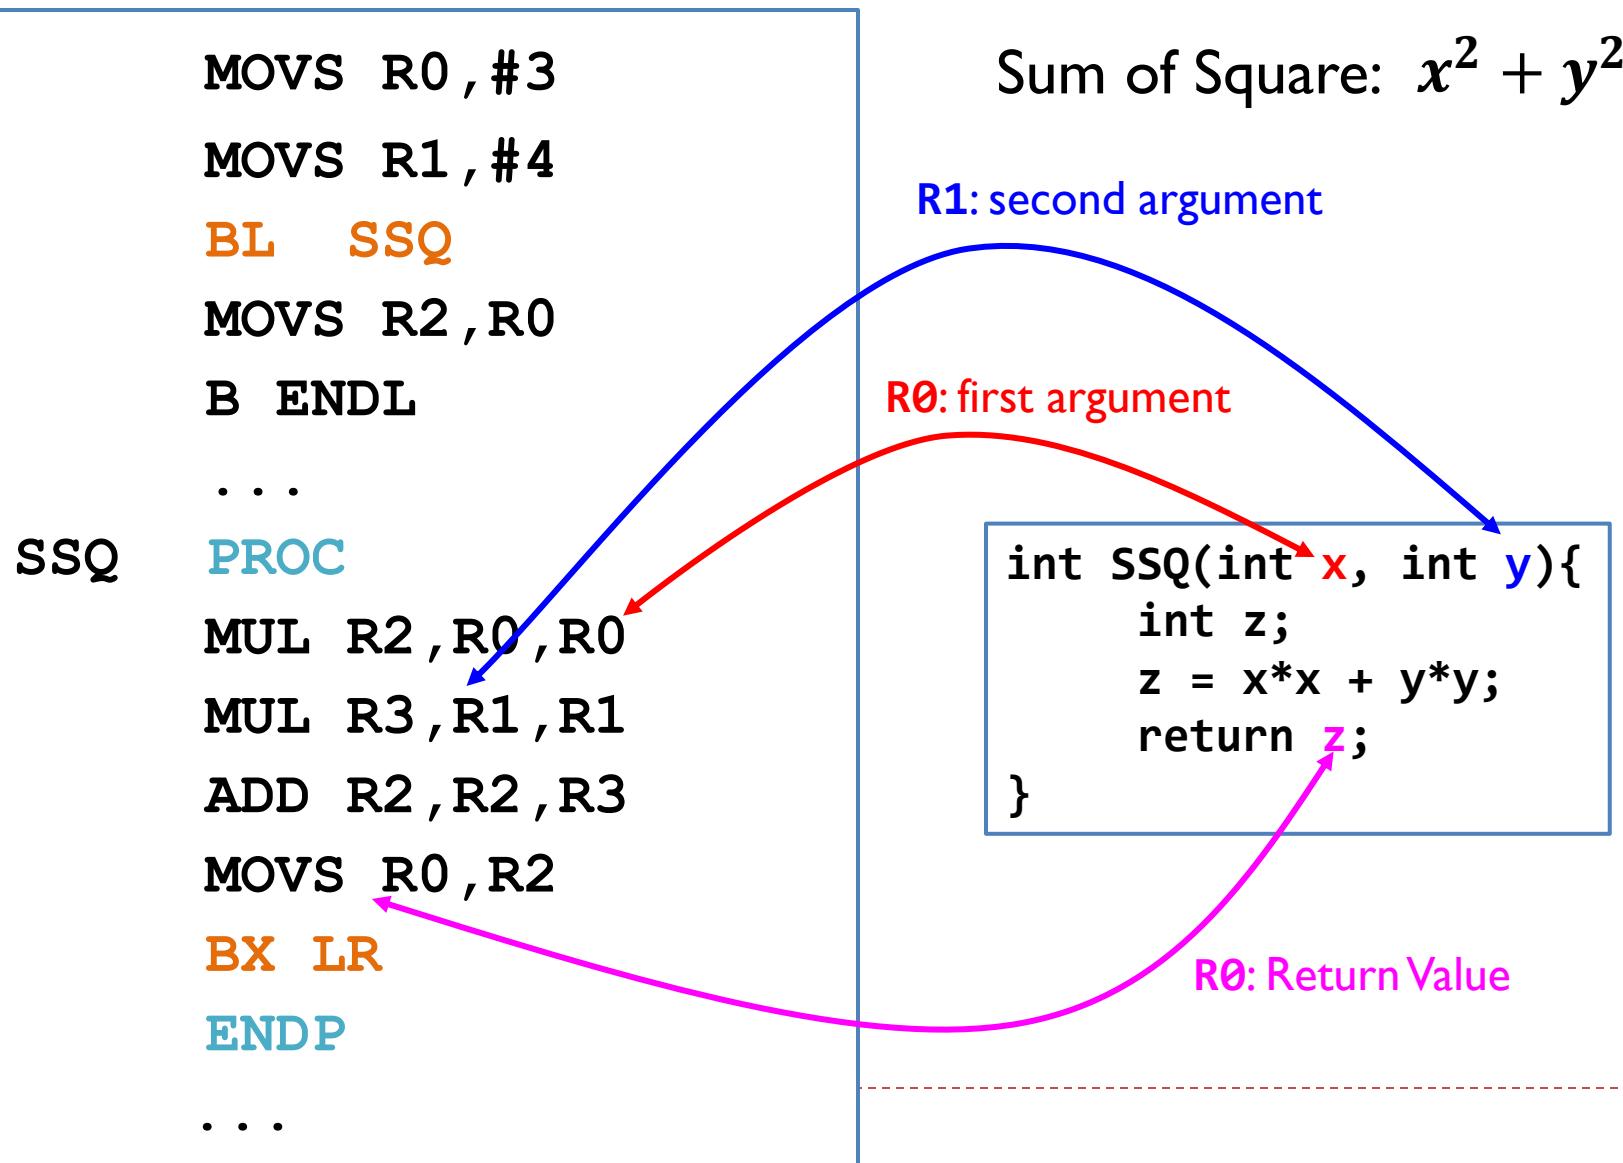
Callee Saved Registers vs Caller Saved Registers



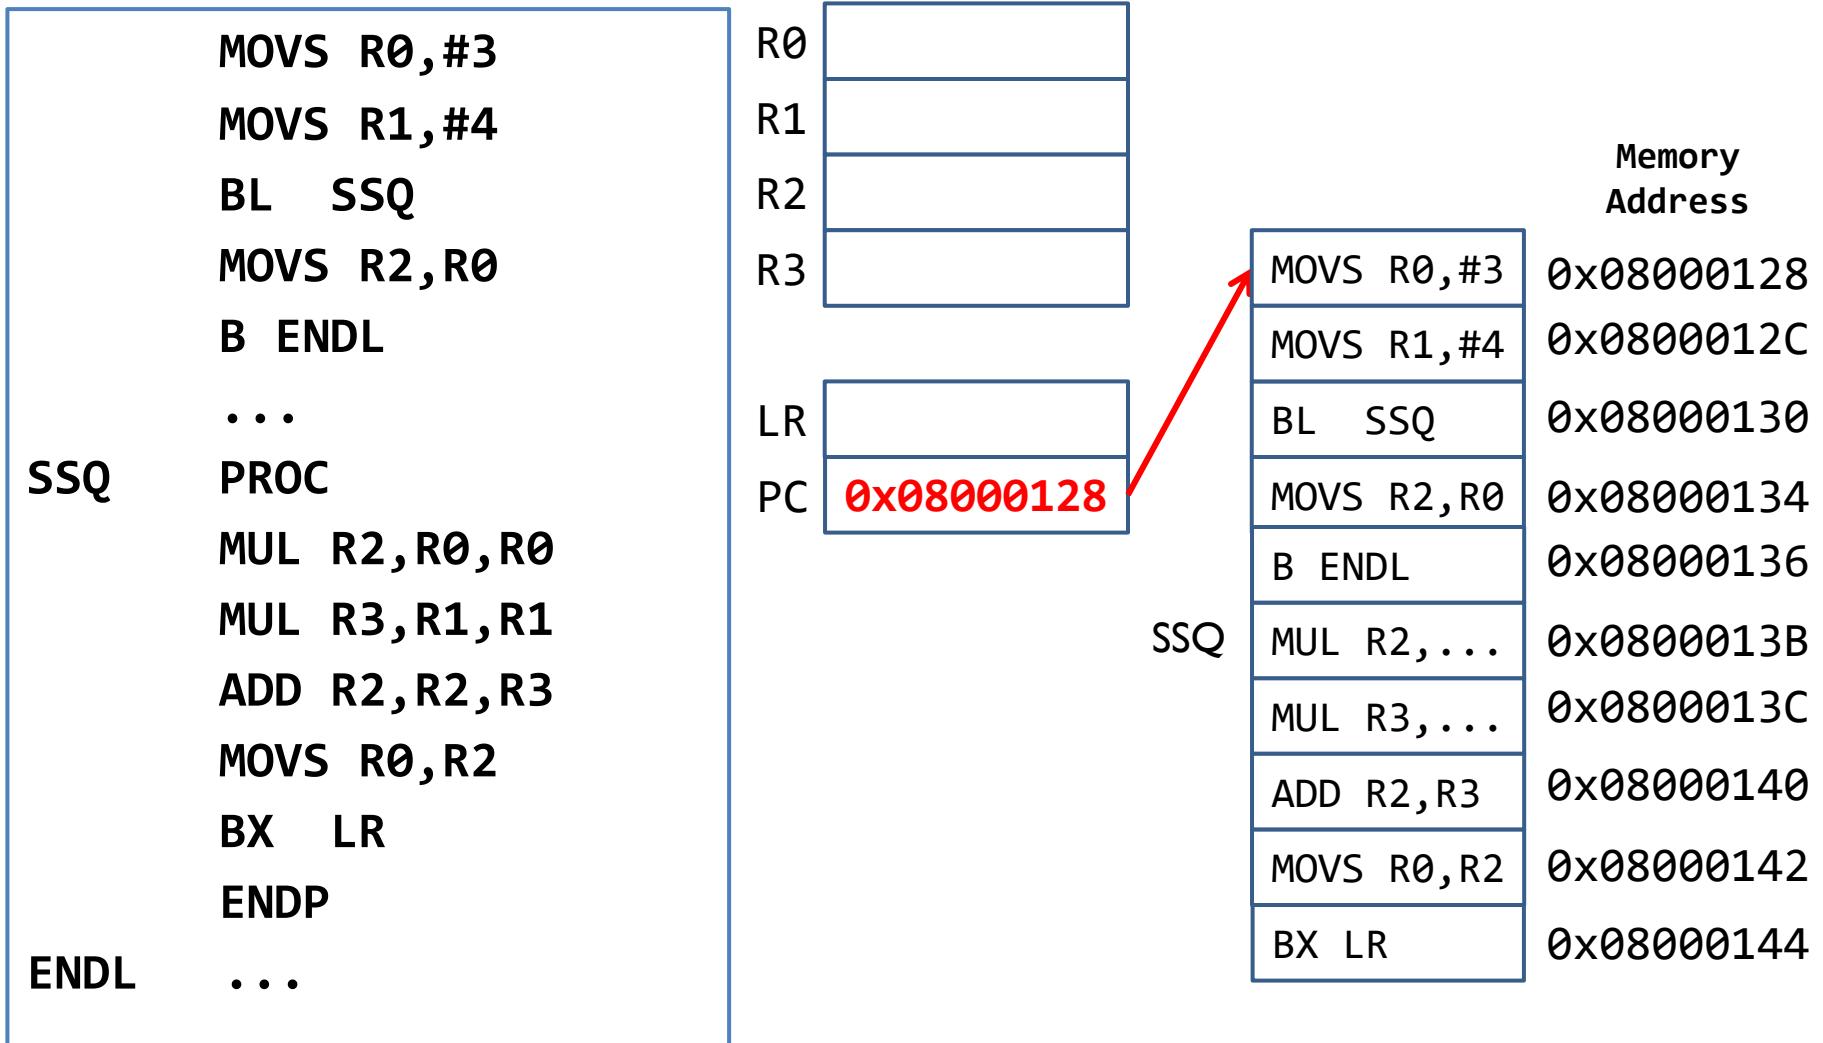
- **Callee can freely modify R0, R1, R2, and R3**
- **If caller expects their values are retained, caller should push them onto the stack before calling the callee**
- **Caller expects these values are retained .**
- **If Callee modifies them, callee must restore their values upon leaving the function.**



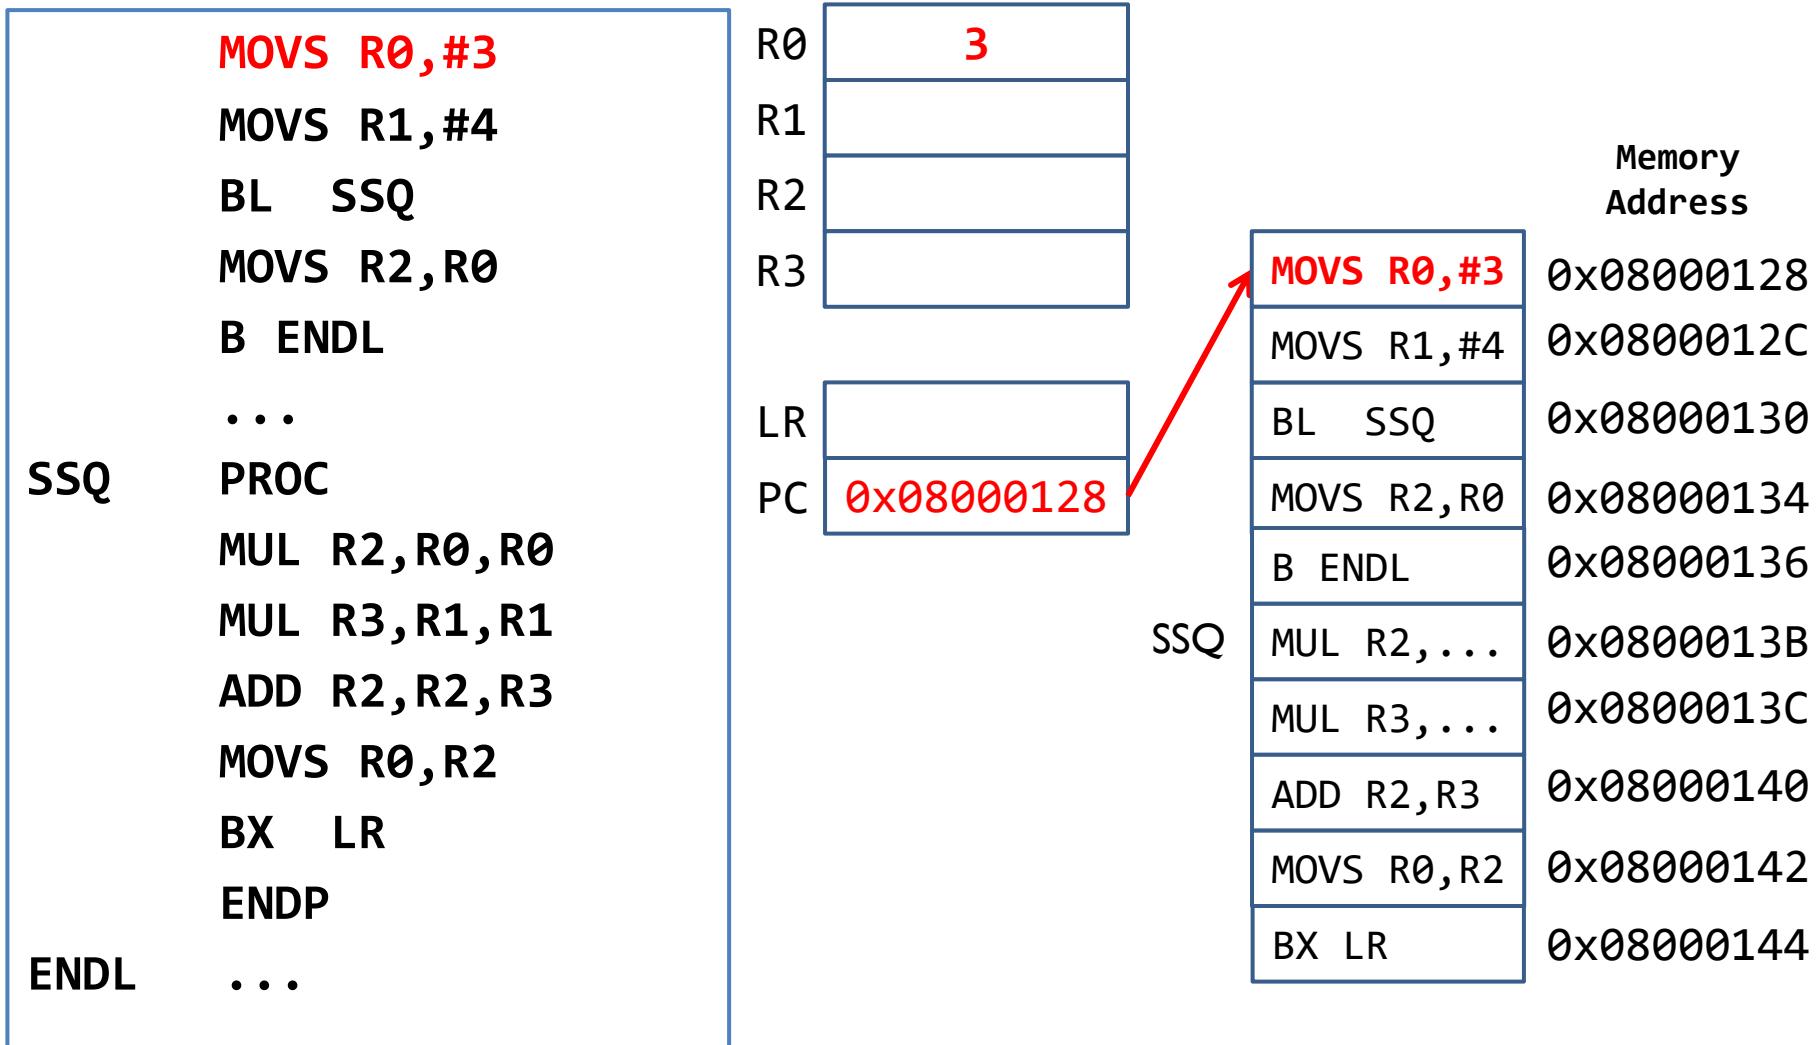
Example: SSQ(3, 4)



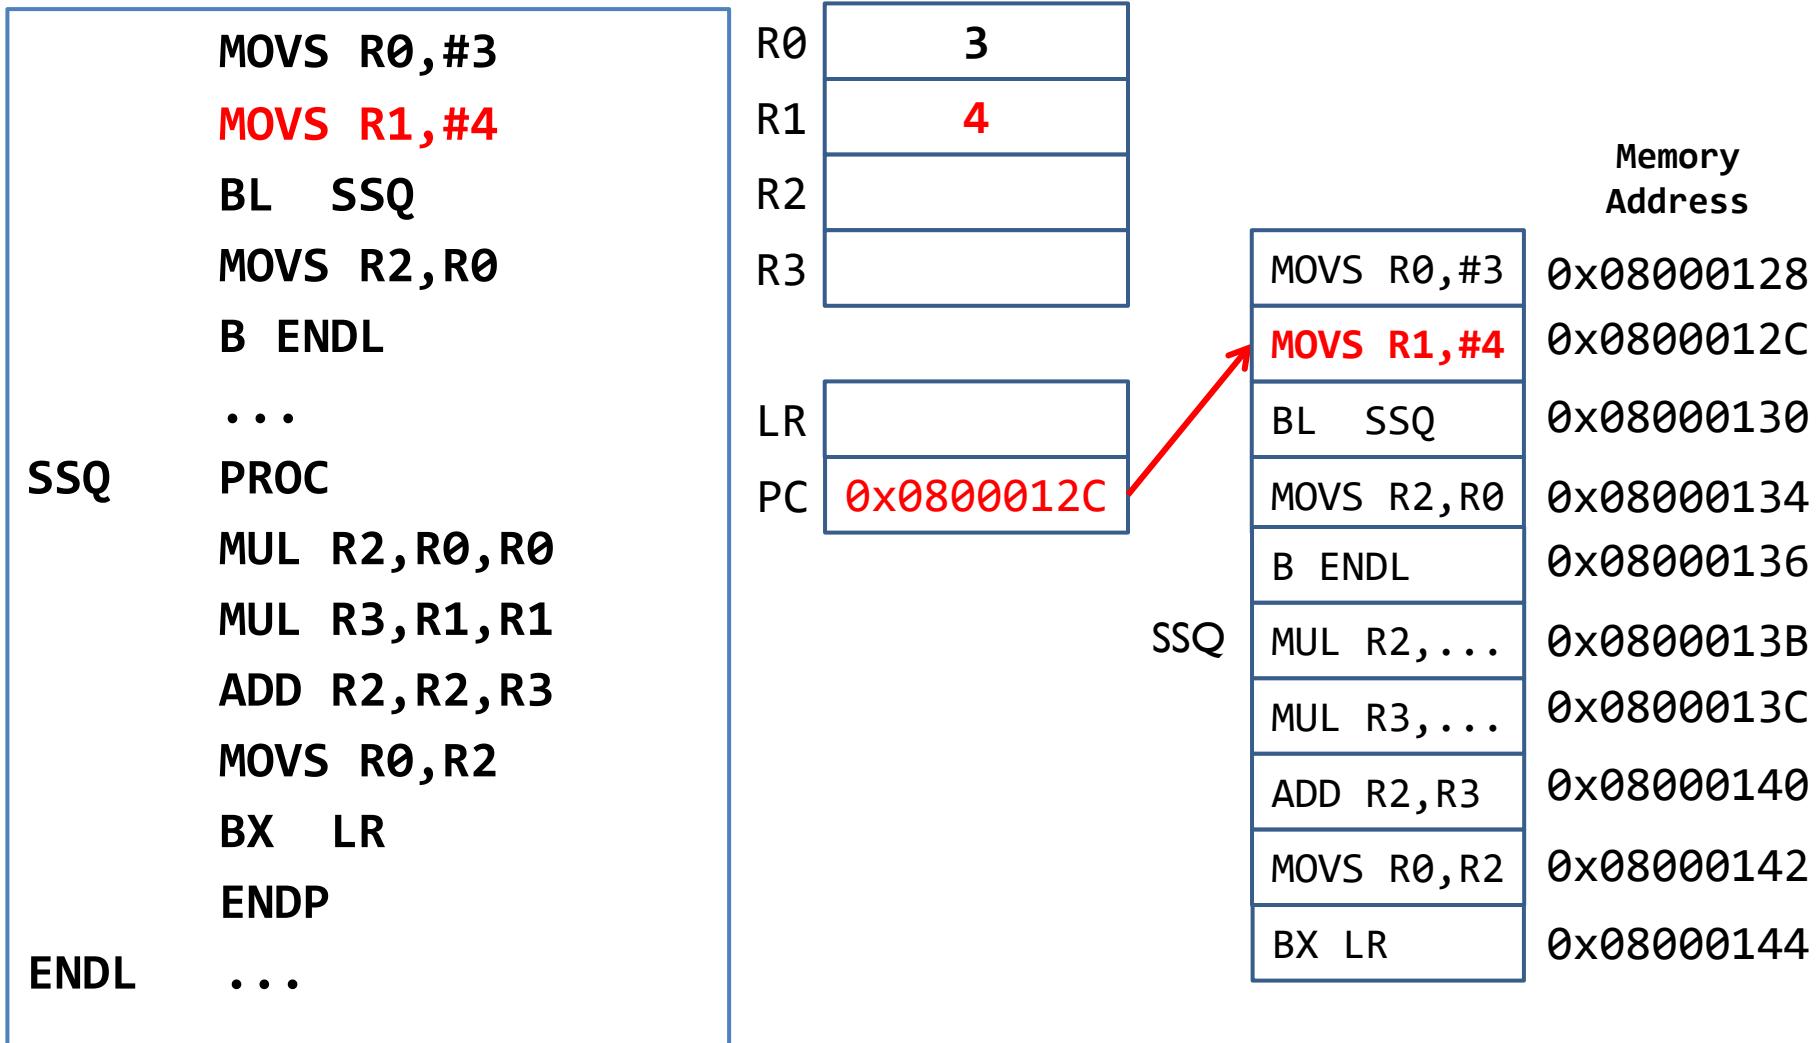
Example: SSQ(3, 4)



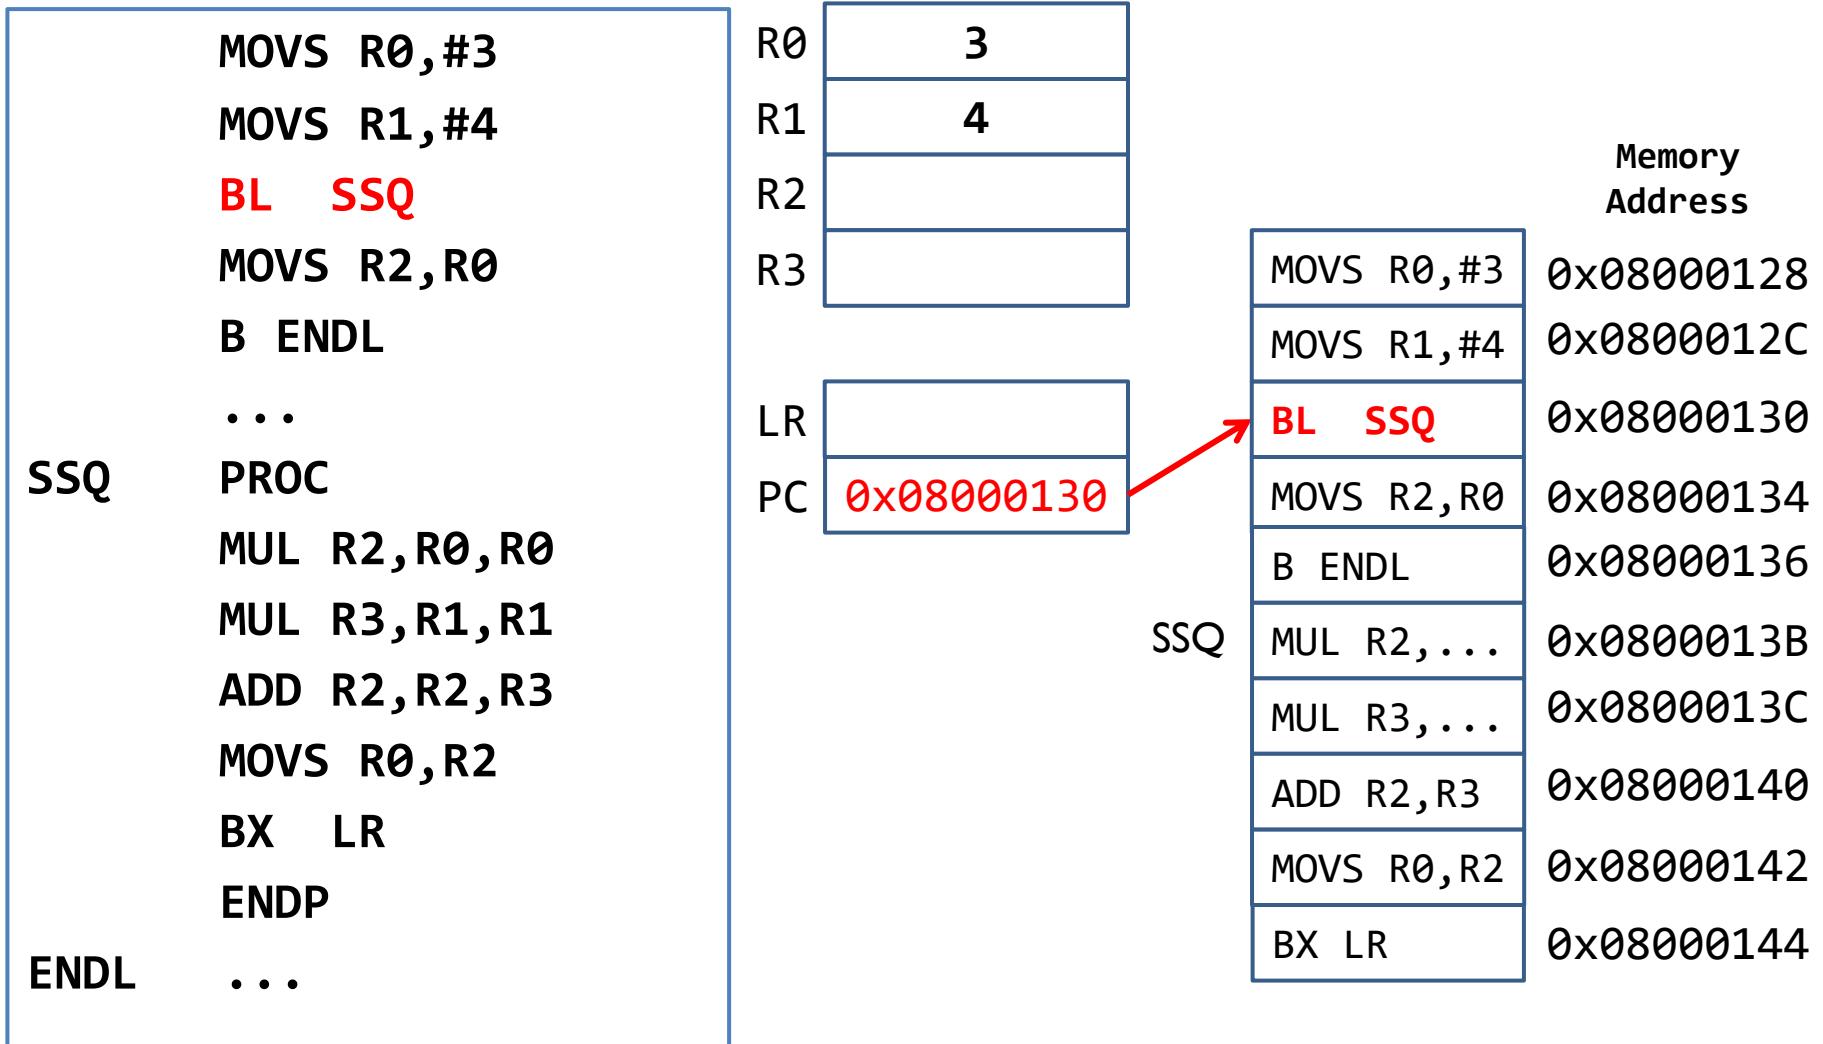
Example: SSQ(3, 4)



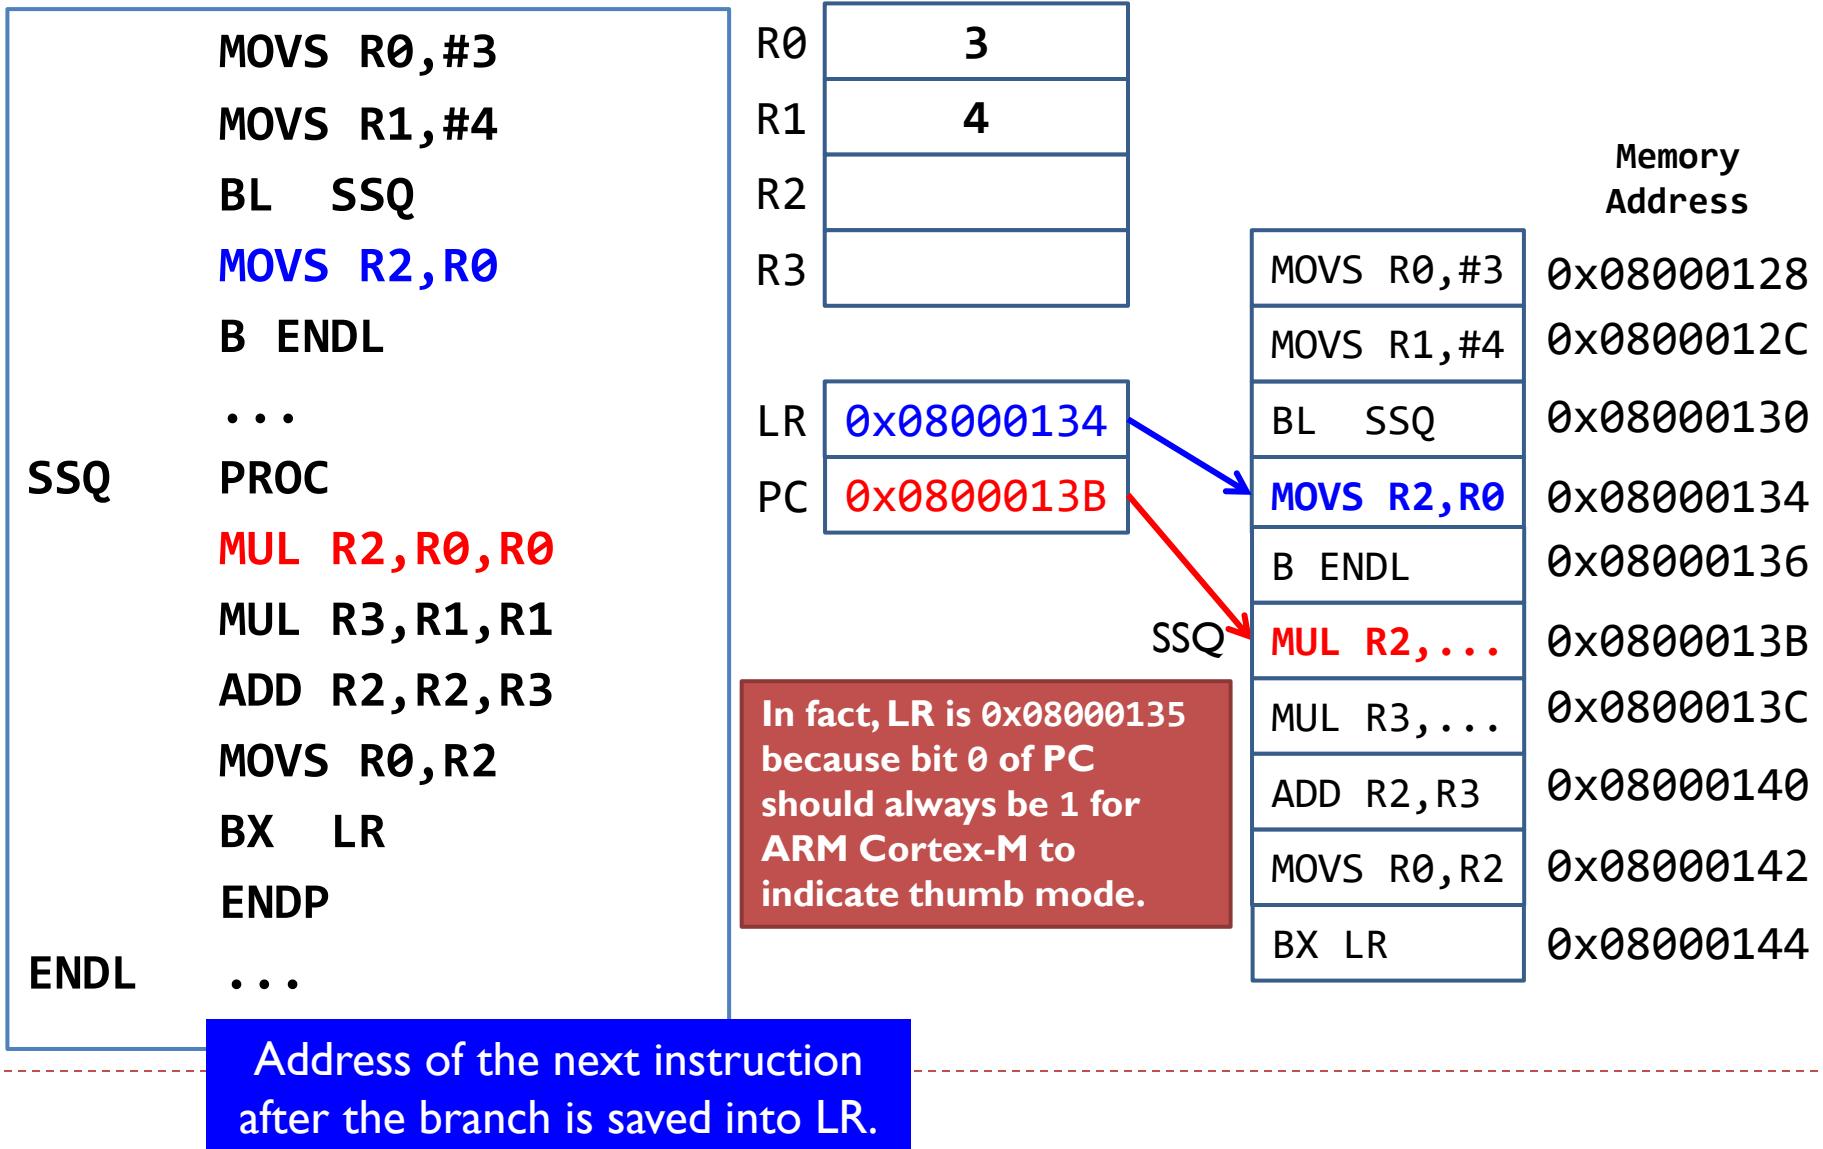
Example: SSQ(3, 4)



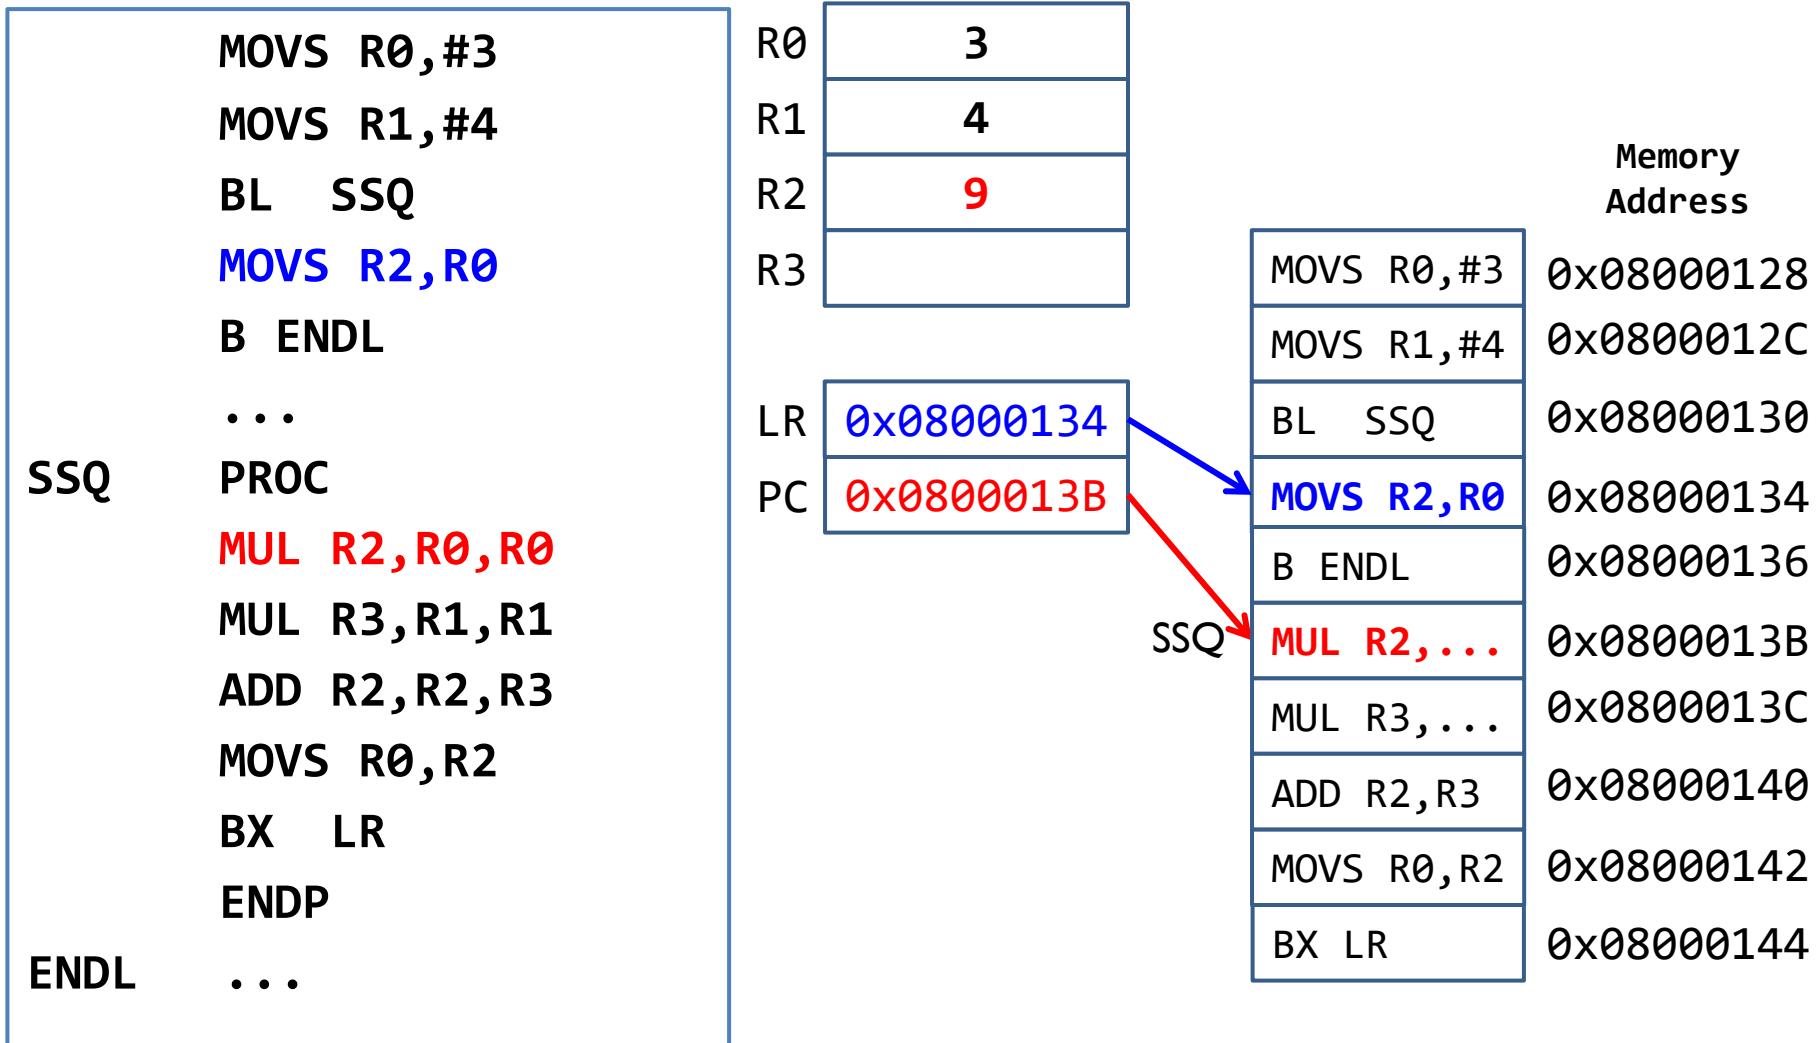
Example: SSQ(3, 4)



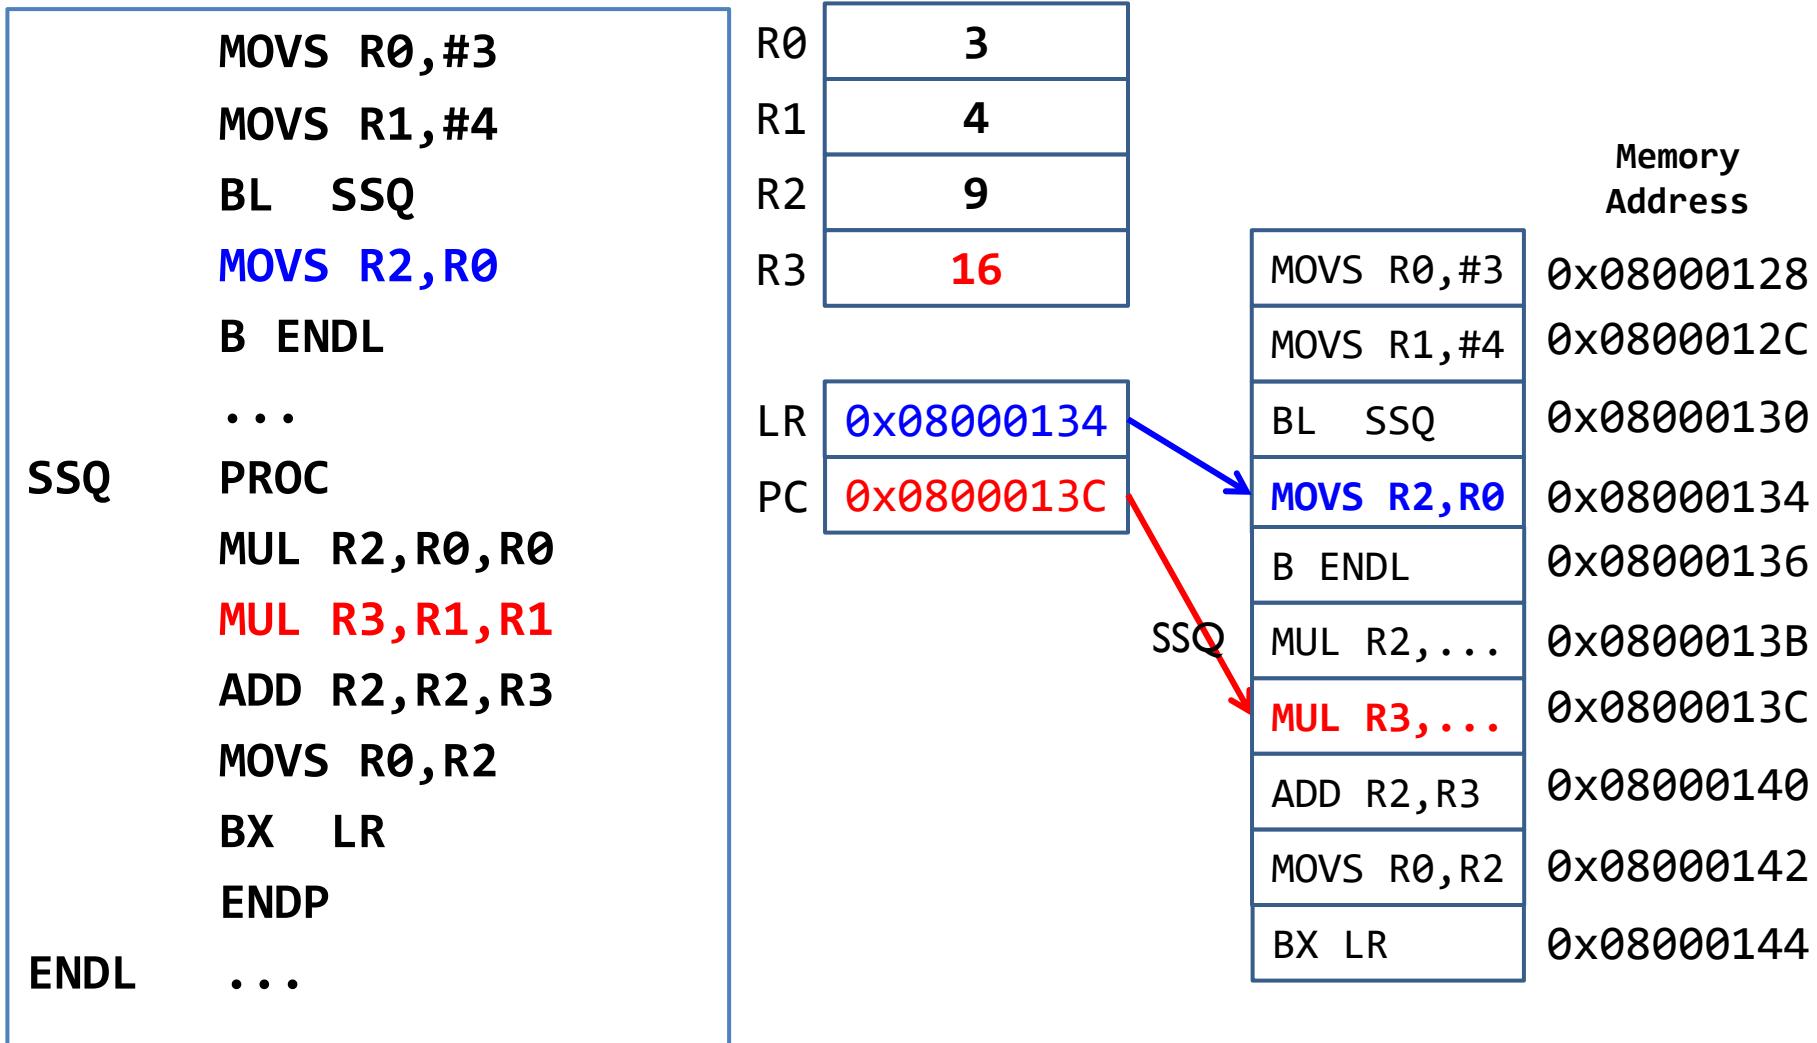
Example: SSQ(3, 4)



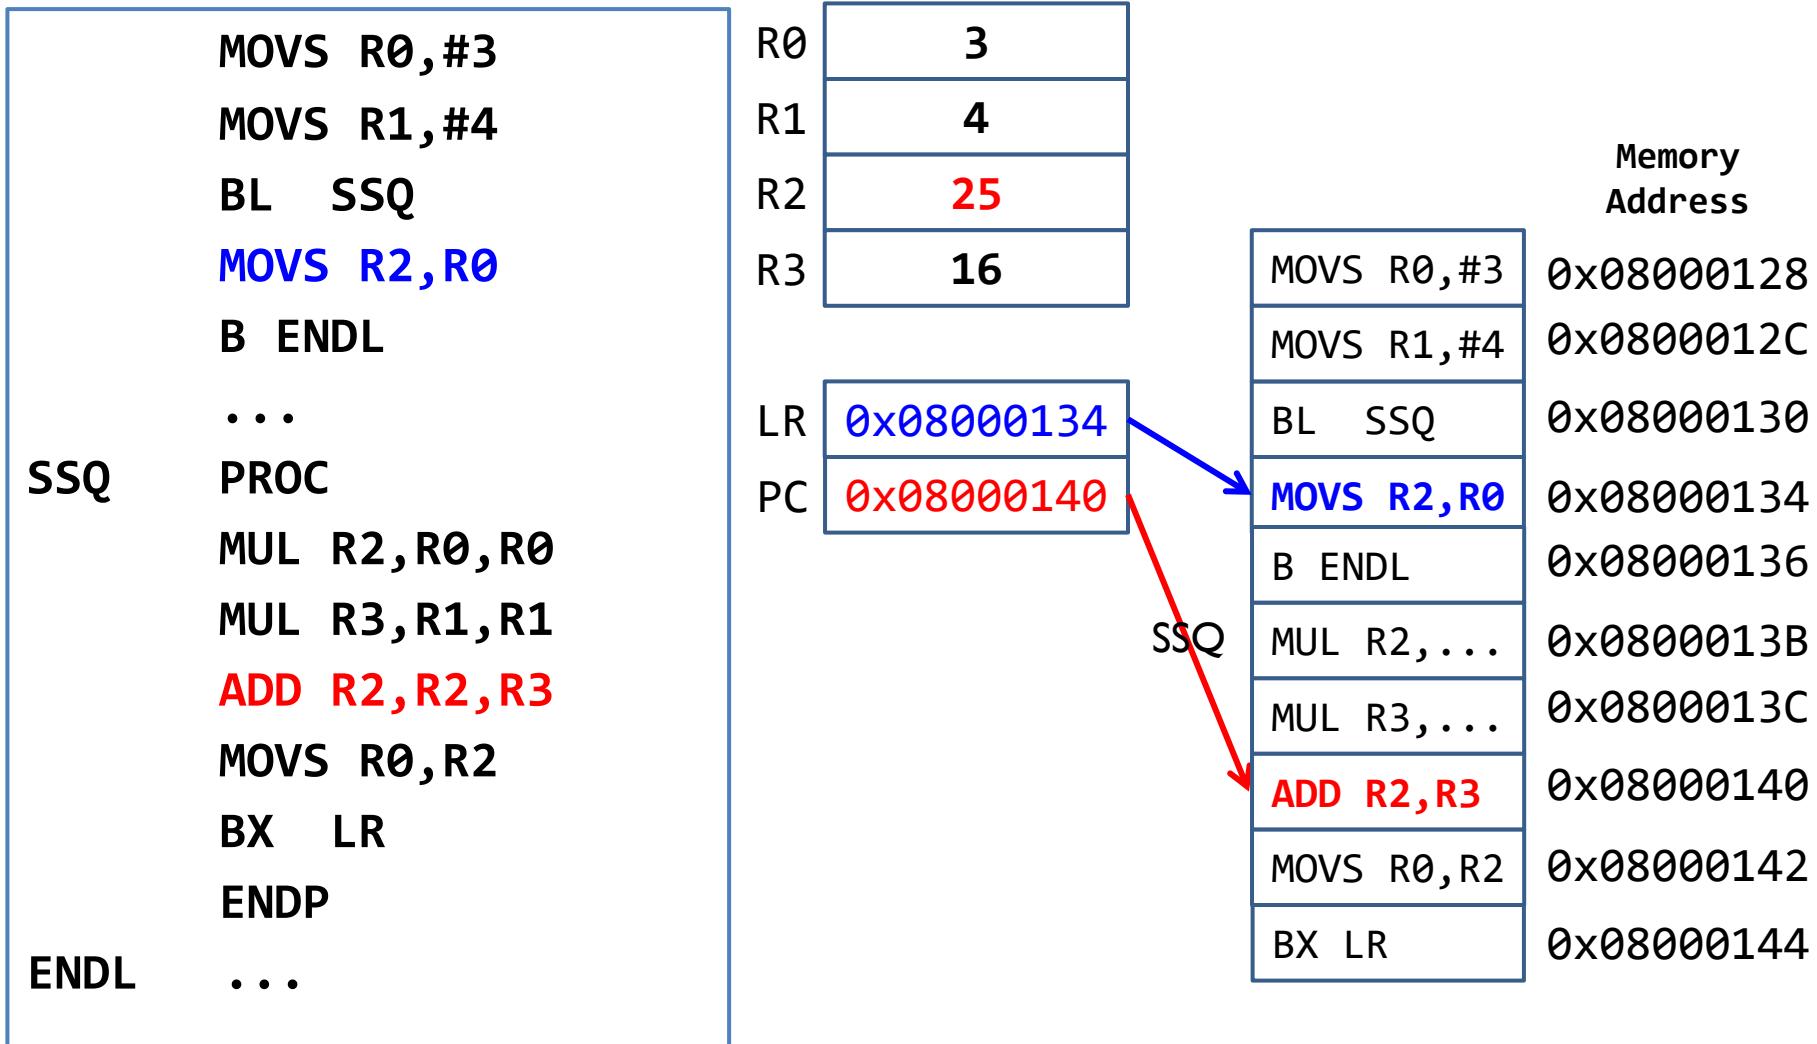
Example: SSQ(3, 4)



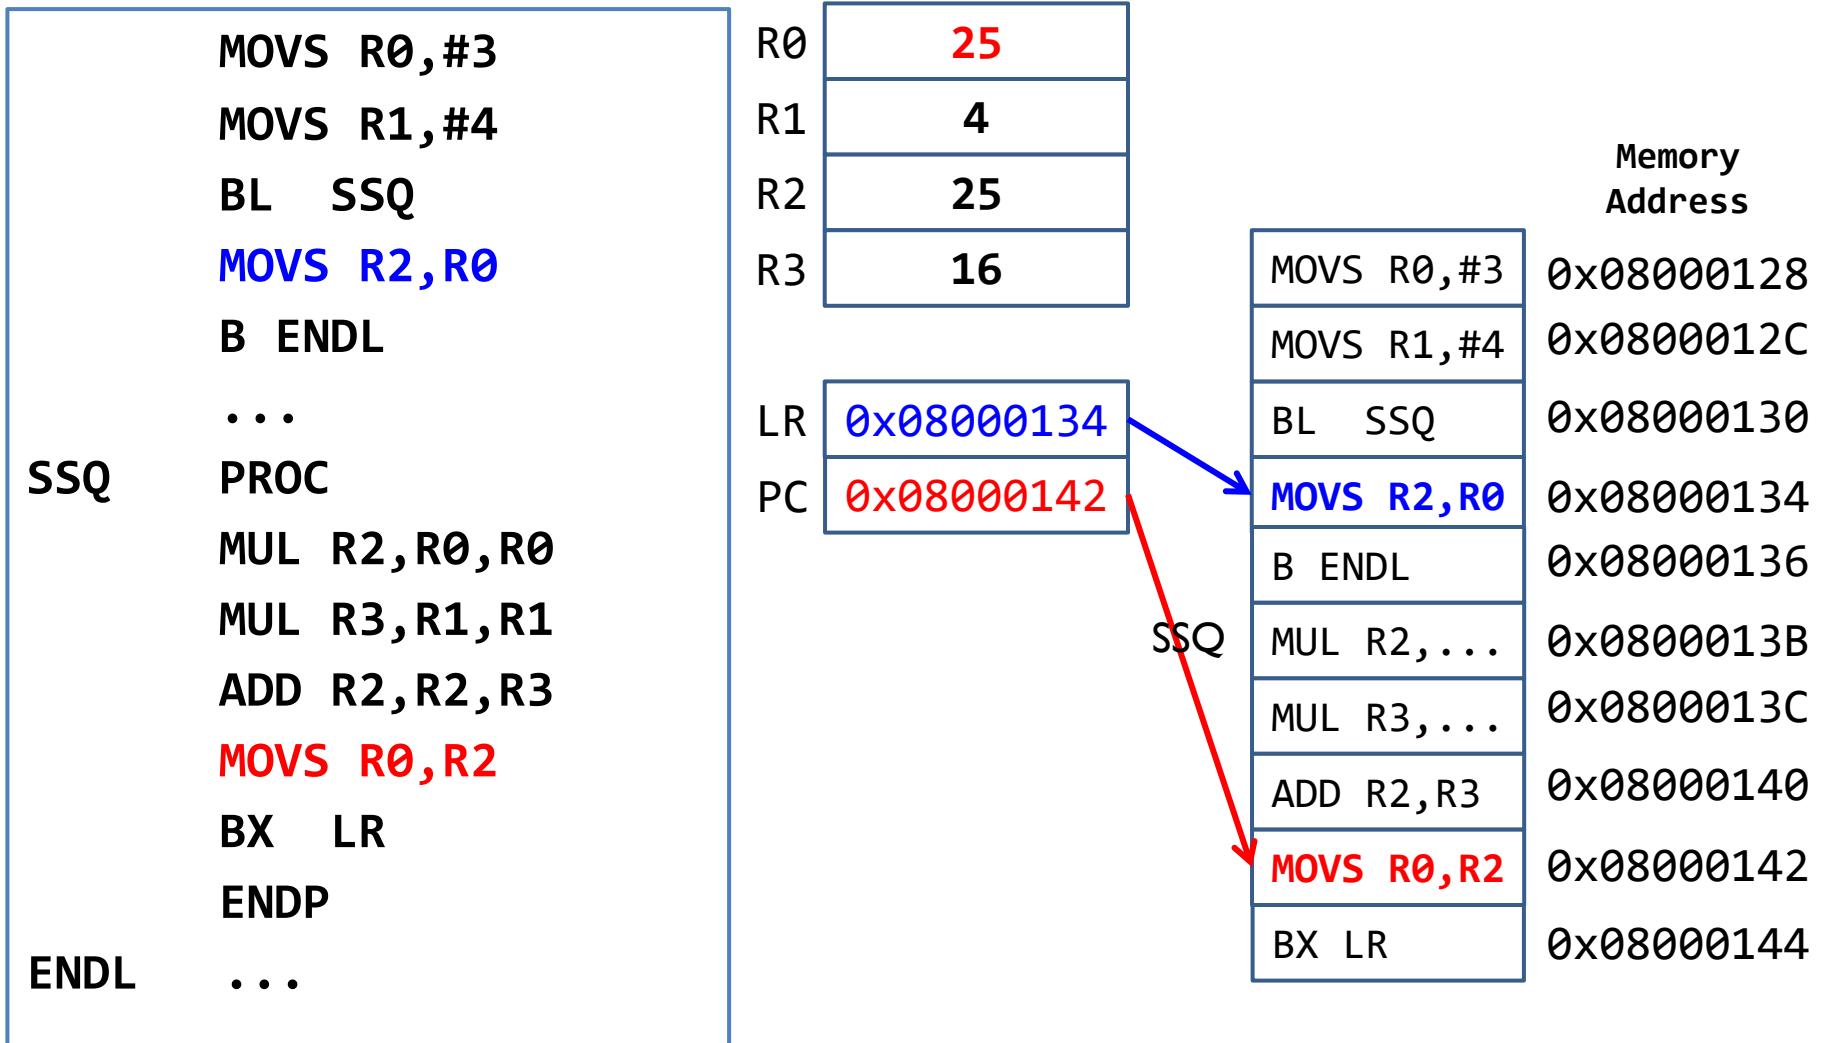
Example: SSQ(3, 4)



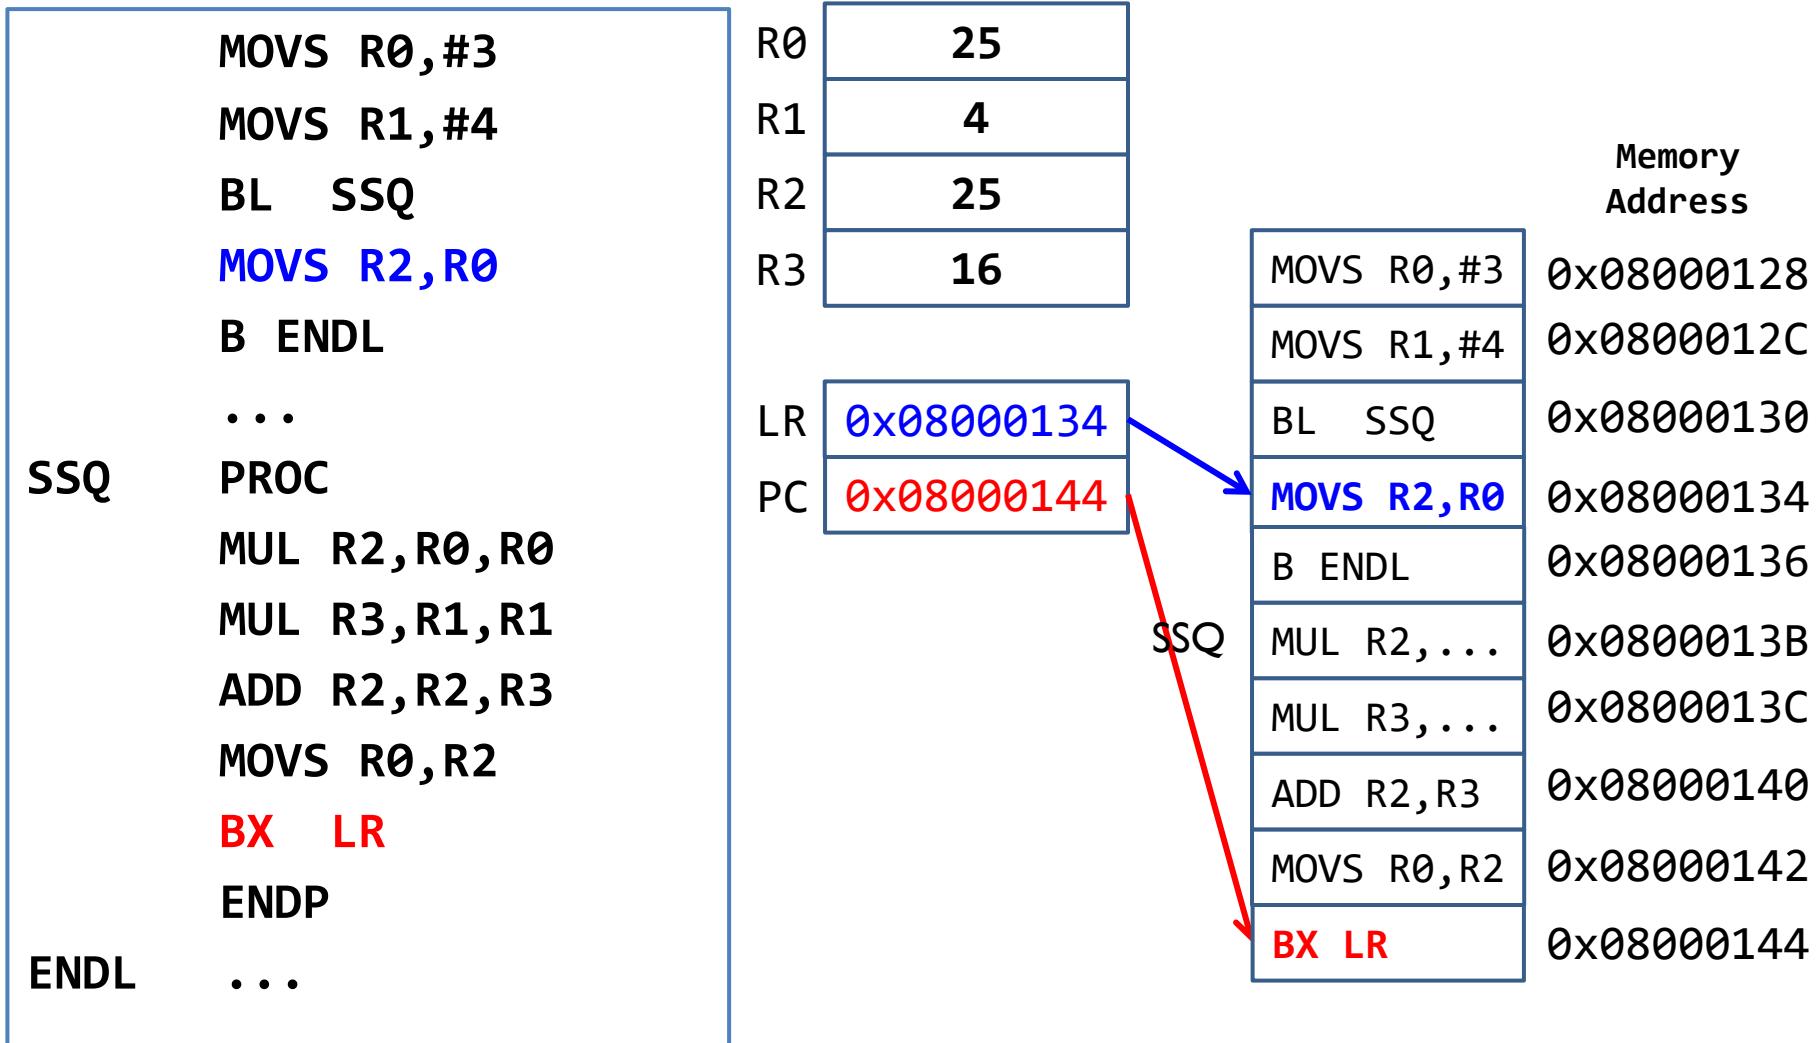
Example: SSQ(3, 4)



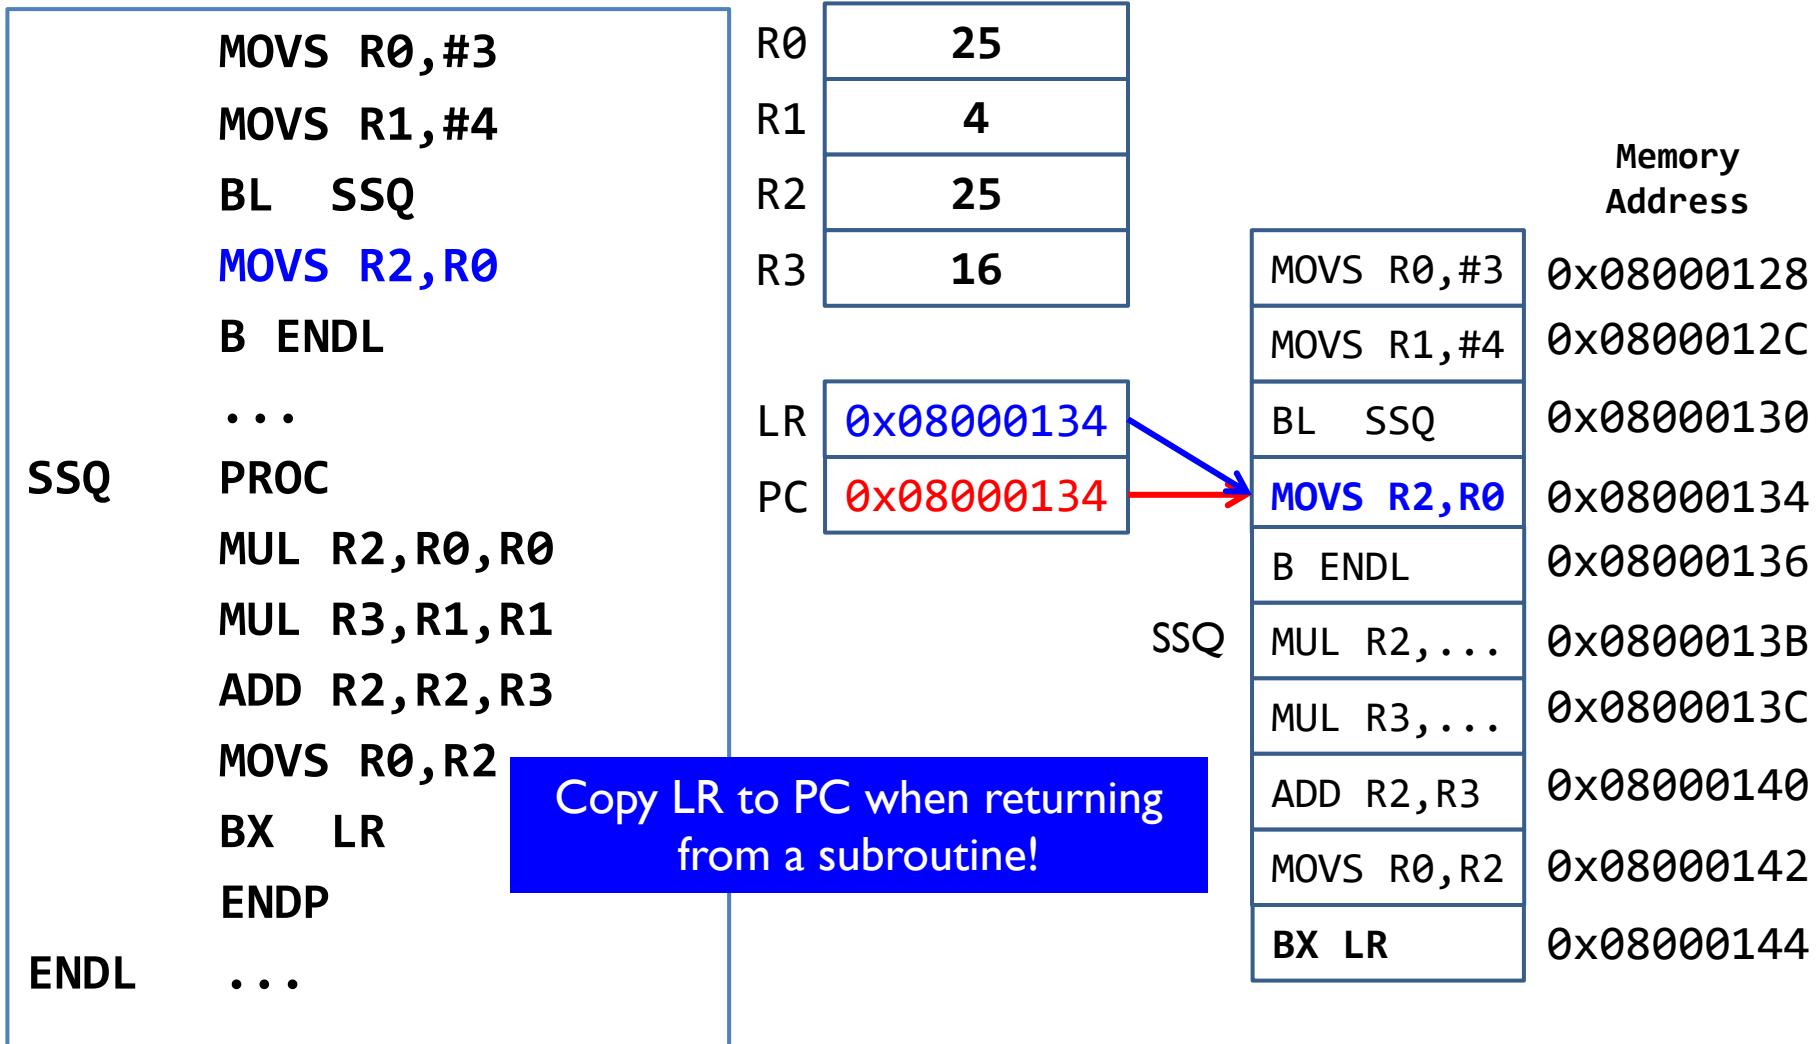
Example: SSQ(3, 4)



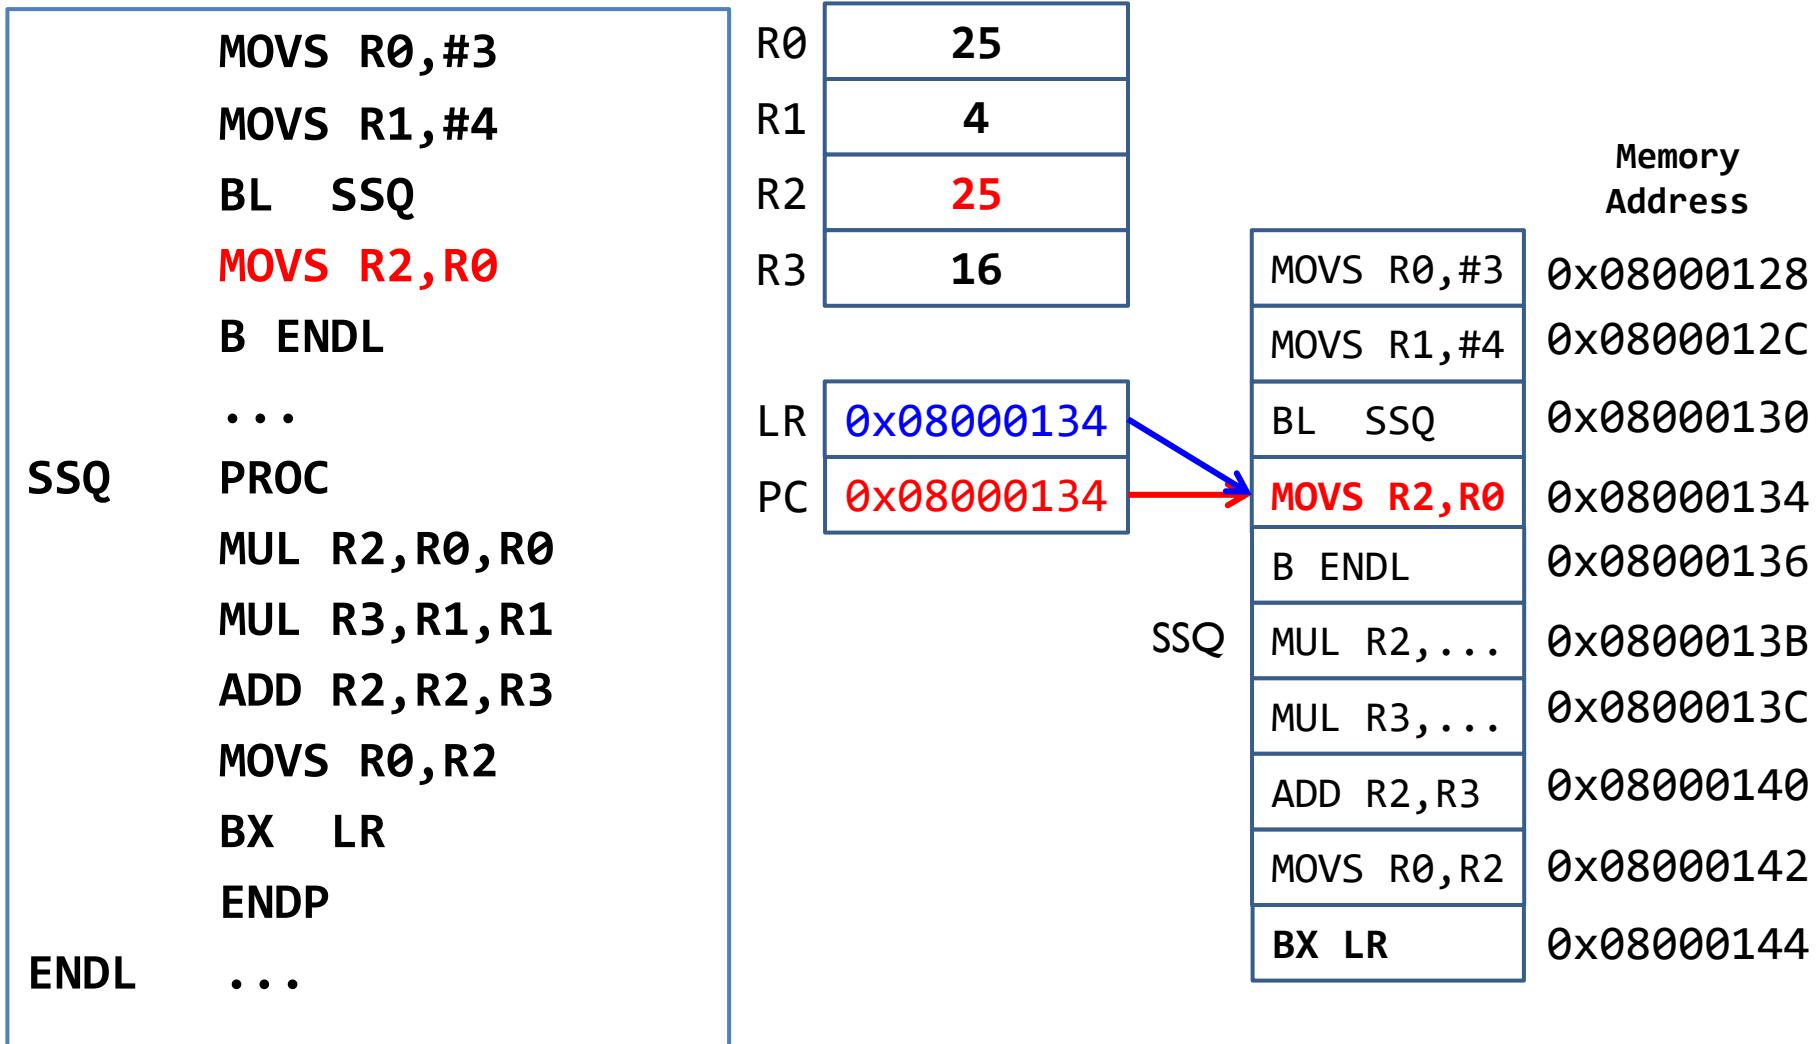
Example: SSQ(3, 4)



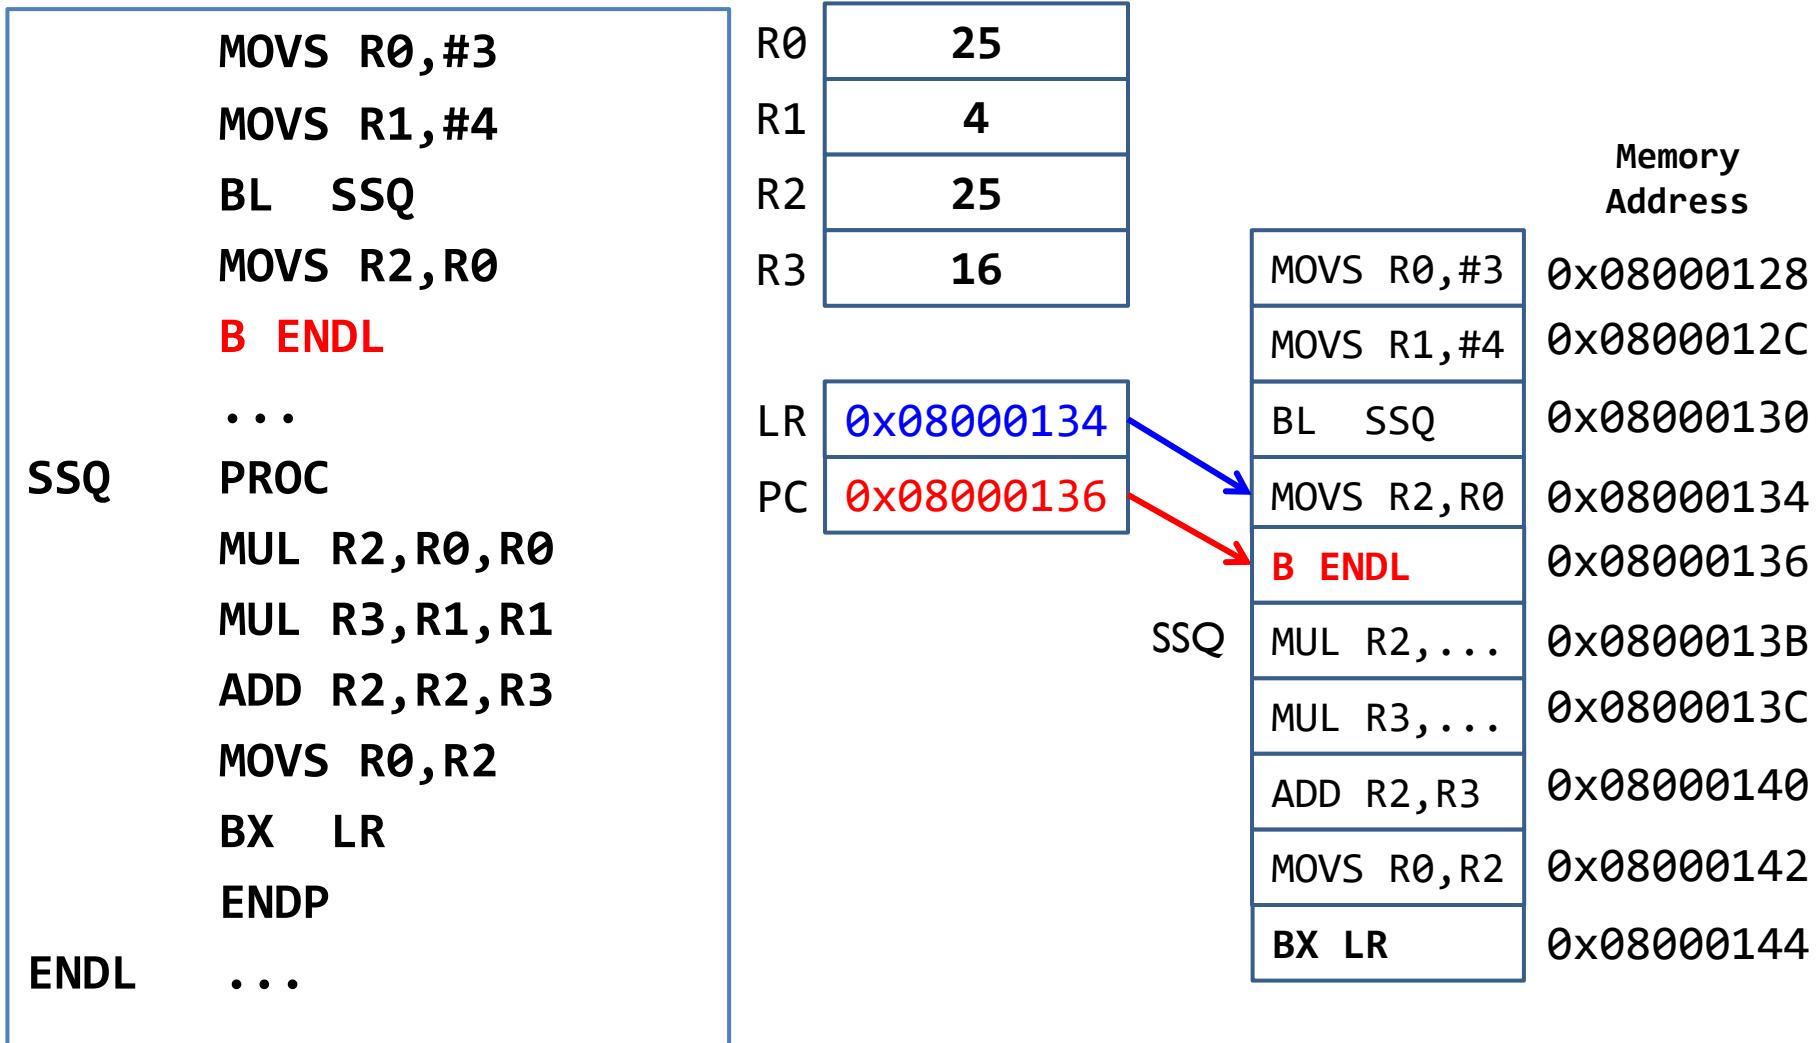
Example: SSQ(3, 4)



Example: SSQ(3, 4)



Example: SSQ(3, 4)



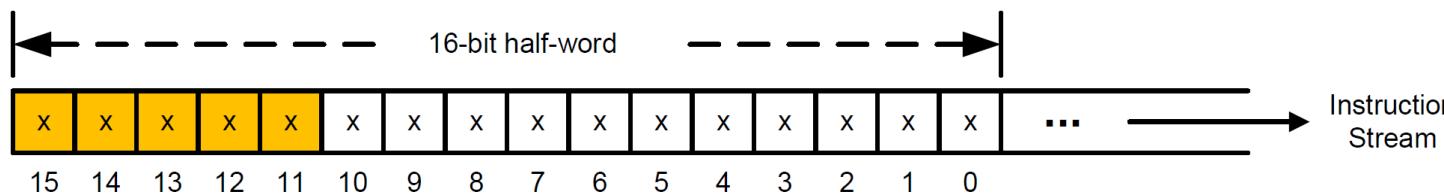
Realities

- ▶ In the previous example,
 - ▶ PC is incremented by 2 or 4.
 - ▶ The least significant bit of LR is always 0.

Well, I lied!

Realities

- ▶ PC is always incremented by **4**.
 - ▶ Each time, 4 bytes are fetched from the instruction memory
 - ▶ It is either two 16-bit instructions or one 32-bit instruction



If bit [15-11] = **11101**, **11110**, or **11111**, then, it is the first half-word of a 32-bit instruction. Otherwise, it is a 16-bit instruction.

- ▶ The least significant bit of LR is always **1** for ARM Cortex-M
 - ▶ This bit is used to control the processor mode:
 - ▶ 0 = ARM, 1 = THUMB
 - ▶ Cortex-M only supports THUMB.

Summary

- ▶ How to call a subroutine?
 - ▶ Branch with link: **BL subroutine**
- ▶ How to return the control back to the caller?
 - ▶ Branch and exchange: **BX LR**
- ▶ How to pass arguments into a subroutine?
 - ▶ Each 8-, 16- or 32-bit variables is passed via r0, r1, r2, r3
 - ▶ Extra parameters are passed via stack
- ▶ How to return a value in a subroutine?
 - ▶ Value is returned in r0
- ▶ How to preserve the running environment for the caller?
 - ▶ (to be covered)

Common Coding Patterns

- ▶ Callee returns a constant in r0.
 - ▶ `mov r0,#17` @ r0 is return value register
 - ▶ `bx lr` @ return from function
- ▶ Callee saves some registers, does some arithmetic, and returns the result in r0.
 - ▶ `push {r4-r7,lr}`
 - ▶ `mov r4,#10`
 - ▶ `mov r5,#100`
 - ▶ `add r0,r4,r5`
 - ▶ `pop {r4-r7,pc}` @ pop saved lr value into PC to return from function
- ▶ Callee calls another function (nested function calls)
 - ▶ `push {lr}` @ must save LR if we call our own function
 - ▶ `mov r0,#123` @ r0 is first function parameter
 - ▶ `bl print_int` @ call function `print_int(123)`
 - ▶ `pop {pc}` @ pop saved lr into PC to return from function
- ▶ Callee return: restore previously-pushed LR, then jump to LR (POP {lr}; BX lr), or equivalently, pop previously-pushed LR to PC
 - ▶ `POP {pc} ≡ POP {lr}; BX lr`

Common Coding Patterns

- ▶ Memory access: first put memory address into register, then load memory content at that address
 - ▶ `adr r2, mydata` @Compute address of label mydata using a PC-relative add and put that address in r2
 - ▶ `ldr r0, [r2]` @Dereference that address, loading the 32-bit word stored at mydata into r0
 - ▶ `bx lr`
 - ▶ `mydata:`
 - ▶ `.word 123`
- ▶ Or
 - ▶ `ldr r2,=mydata` @ pseudo-instruction that loads absolute address of mydata from a nearby literal pool into r2
 - ▶ `ldr r0,[r2]`
 - ▶ `bx lr`
 - ▶ `mydata:`
 - ▶ `.word 123`
- ▶ adr vs. ldr
 - ▶ If mydata is in range for adr (c.f.,), both forms will leave r2 holding the same address at run time.
 - ▶ Out-of-range labels: adr may fail; ldr =mydata still works, but it incurs one memory access
- ▶ Note: pseudo-instruction LDR Rd, =X loads a value if X is a constant, and an address if X is a symbol,
 - ▶ e.g., “LDR r2, =0x55555555” loads the value 0x55555555 from memory

References

- ▶ **Lecture 29. Calling a subroutine**
 - ▶ <https://www.youtube.com/watch?v=xt2Q9n1Udb4&list=PLRJhV4hUhlymmp5CCeIFPyxbknsdcXCc8&index=29>
- ▶ **Lecture 30. Passing Arguments to a Subroutine**
 - ▶ <https://www.youtube.com/watch?v=DGKjFKjxAYs&list=PLRJhV4hUhlymmp5CCeIFPyxbknsdcXCc8&index=31>